



# Free lunch is over

## WHY C++ IS SO IMPORTANT IN THE MODERN WORLD?

Mateusz Pusz

September 21, 2019

# Confidential Information

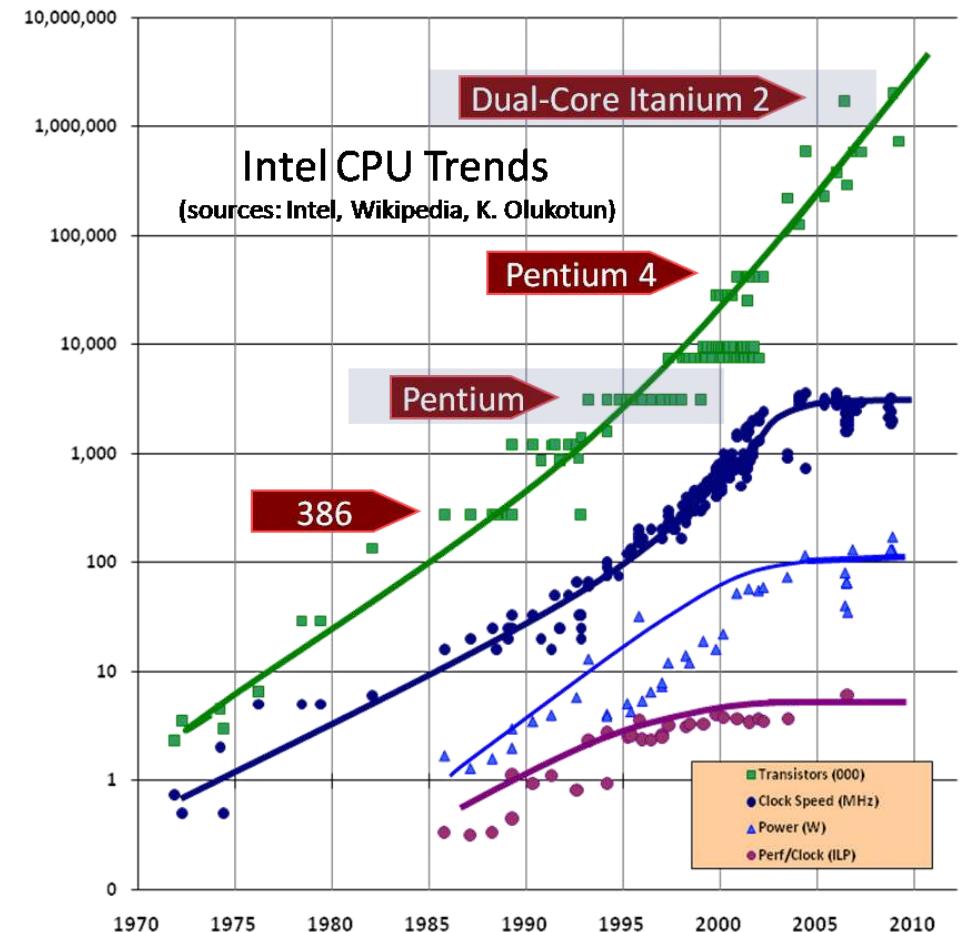
---

This presentation has been prepared by EPAM Systems, Inc. solely for use by EPAM at its EPAM Software Engineering Conference. This presentation or the information contained herein may not be reproduced or used for any other purpose. This presentation includes highly confidential and proprietary information and is delivered on the express condition that such information will not be disclosed to anyone except persons in the recipient organization who have a need to know solely for the purpose described above. No copies of this presentation will be made, and no other distribution will be made, without the consent of EPAM. Any distribution of this presentation to any other person, in whole or in part, or the reproduction of this presentation, or the divulgence of any of its contents is unauthorized.

## **WHY C++?**

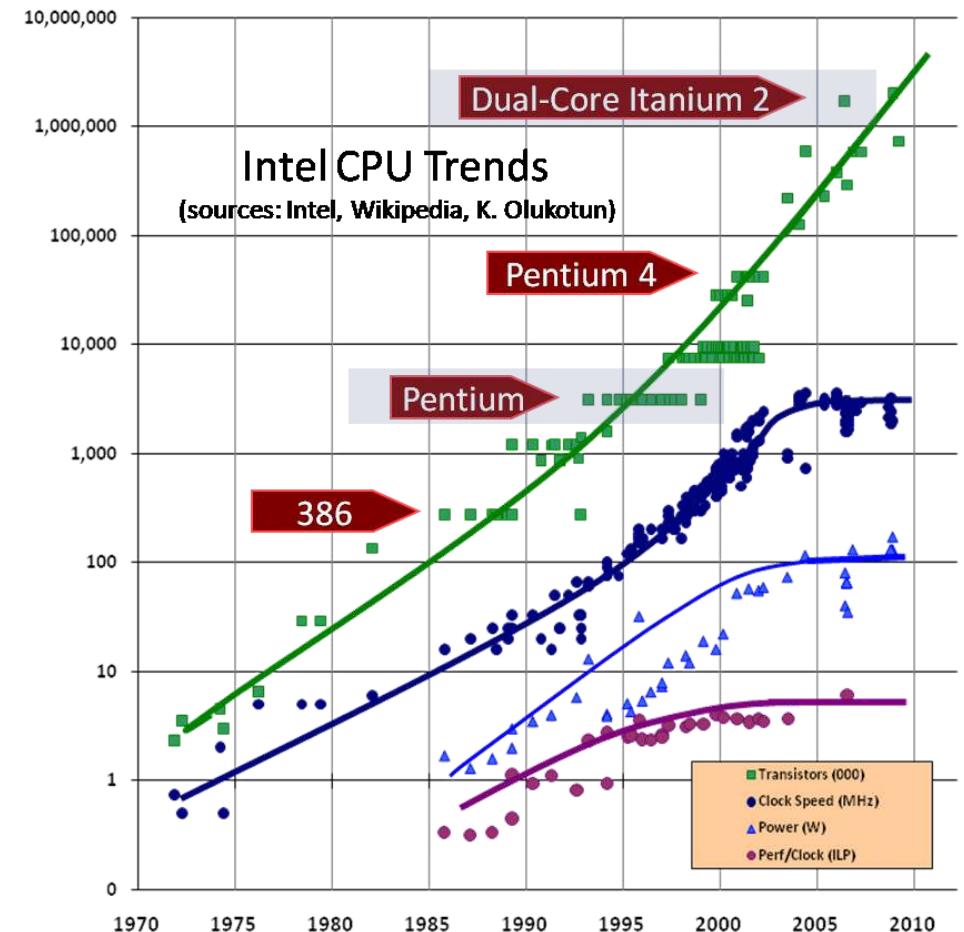
# Herb Sutter 2005: Free Lunch is Over

- No matter how fast processors get, *software consistently finds new ways to eat up the extra speed*



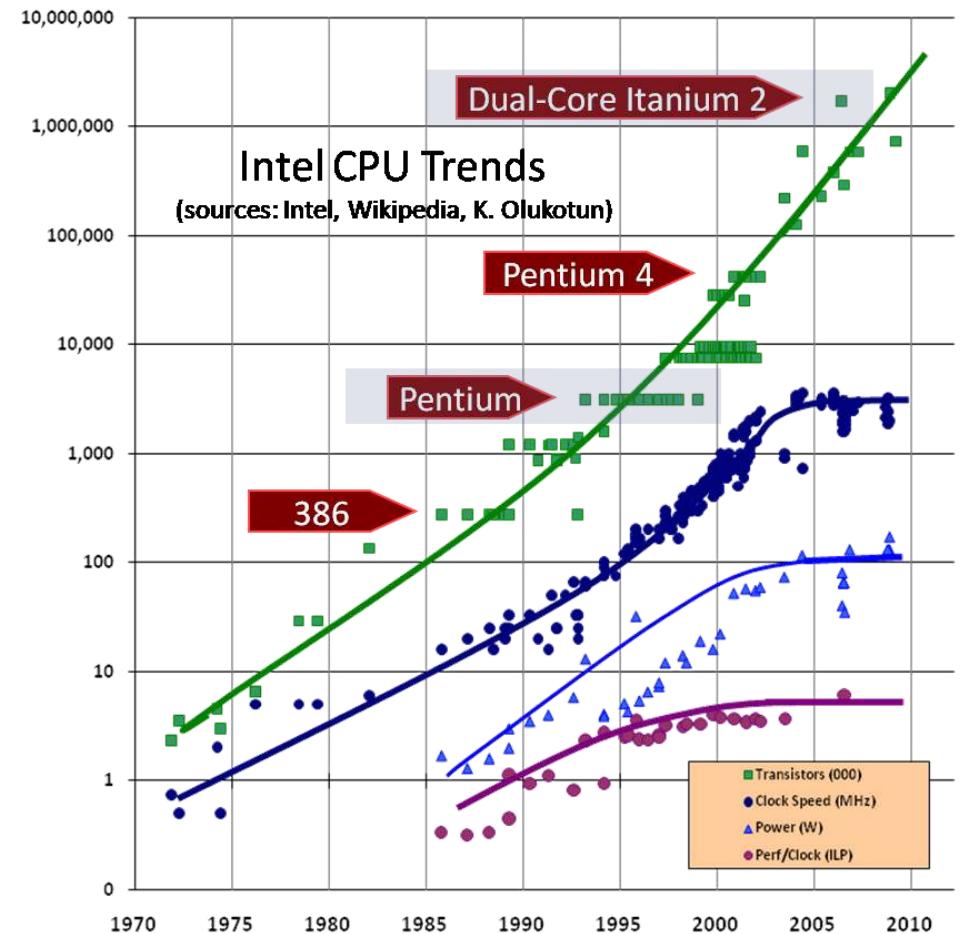
# Herb Sutter 2005: Free Lunch is Over

- No matter how fast processors get, *software consistently finds new ways to eat up the extra speed*
- Applications have enjoyed *free and regular performance gains for several decades*
  - even without releasing new versions



# Herb Sutter 2005: Free Lunch is Over

- No matter how fast processors get, *software consistently finds new ways to eat up the extra speed*
- Applications have enjoyed *free and regular performance gains for several decades*
  - even without releasing new versions
- *Around 2003 we run out of room with approaches to boost CPU clocks*
  - heat (too much and hard to dissipate)
  - power consumption (too high)
  - leakage problems



# Where the Focus Is?

---

LANGUAGE	EFFICIENCY	FLEXIBILITY	ABSTRACTION	PRODUCTIVITY
C	YES!	YES!	non-goal	non-goal
C++	YES!	YES!	YES!	<i>WIP</i>
JAVA, C#	at the expense of	at the expense of	YES!	YES!

# Where the Focus Is?

---

LANGUAGE	EFFICIENCY	FLEXIBILITY	ABSTRACTION	PRODUCTIVITY
C	YES!	YES!	non-goal	non-goal
C++	YES!	YES!	YES!	WIP
JAVA, C#	at the expense of	at the expense of	YES!	YES!

## C++ motto:

You don't pay for what you don't use

# Why C++?

---



# Why C++?

---



## PERFORMANCE PER TRANSISTOR

- Hardware costs
- Limited resources on mobiles

# Why C++?

---



## PERFORMANCE PER TRANSISTOR

- Hardware costs
- Limited resources on mobiles

## PERFORMANCE PER CLOCK

- Get the most from every CPU core
- Bigger experiences on a smaller hardware

# Why C++?

---

PERFORMANCE PER WAT



# Why C++?

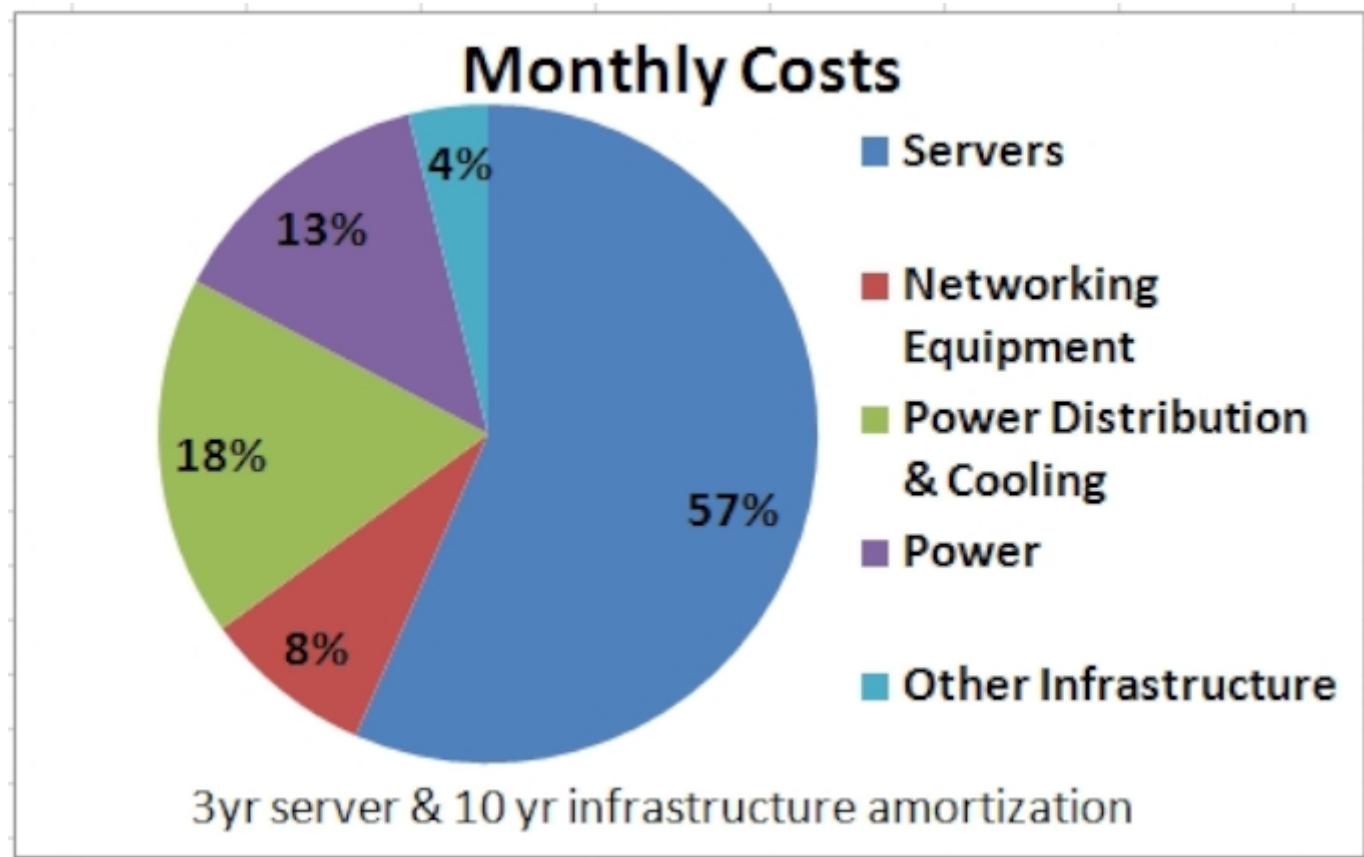
---

## PERFORMANCE PER WAT



# Data center costs

---



<https://perspectives.mvdirona.com/2010/09/overall-data-center-costs>

# How C++ Combats Global Warming

---



# How C++ Combats Global Warming

---

I guess in some way of saying here's my contribution to dealing with global warming because if you can double the efficiency of those systems you don't need yet another server farm that use as much energy as a small town.

-- Bjarne Stroustrup

# Focus shift

---

People costs shift from top to nearly irrelevant.

James Hamilton

VP & Distinguished Engineer, Amazon Web Services

# Efficiency

---

Efficiency is not just running fast or running bigger programs, it's also running using less resources.

-- Bjarne Stroustrup

# Efficiency

---

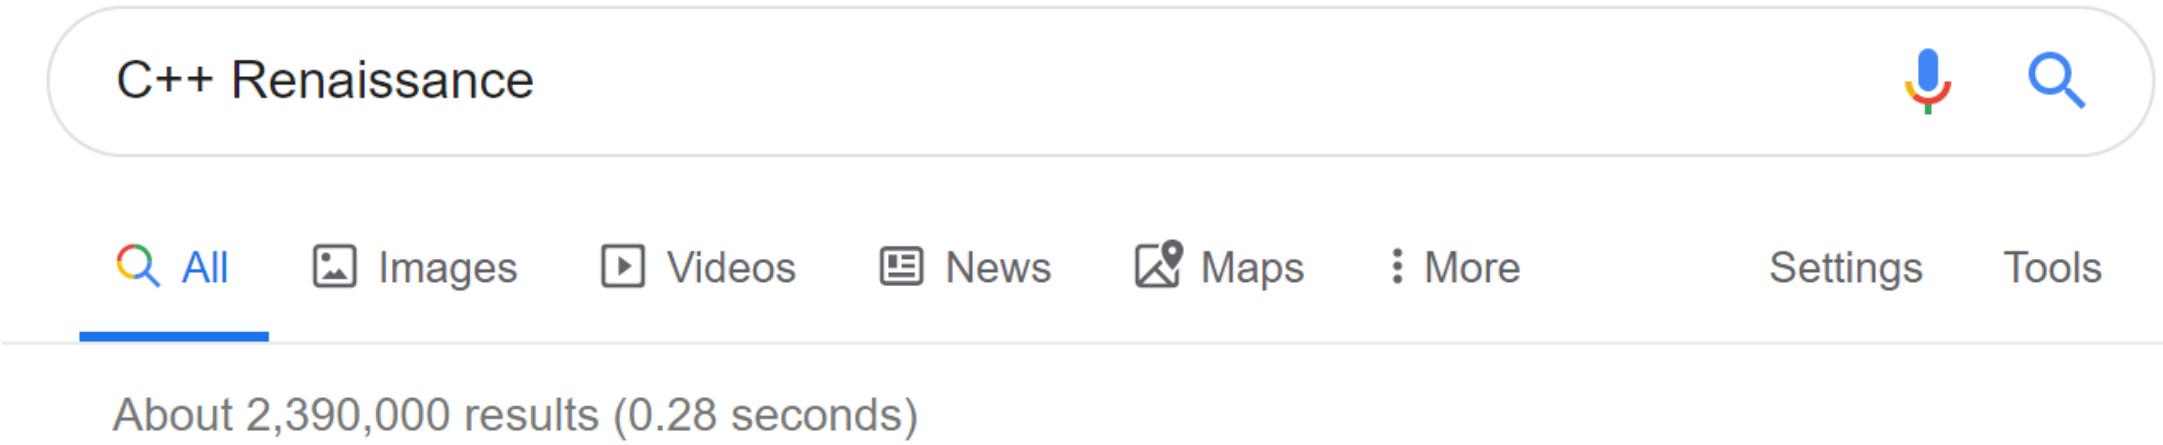
Efficiency is not just running fast or running bigger programs, it's also running using less resources.

-- Bjarne Stroustrup

Performance-oriented corporations are willing to pay more for software development if due to its better efficiency it will save lots of money during product lifetimes

# C++ Renaissance

---



A screenshot of a Google search results page. The search bar at the top contains the query "C++ Renaissance". To the right of the search bar are a microphone icon and a magnifying glass icon. Below the search bar is a navigation bar with categories: All (selected), Images, Videos, News, Maps, More, Settings, and Tools. A blue horizontal bar is under the "All" category. Below the navigation bar, the text "About 2,390,000 results (0.28 seconds)" is displayed.

Monday, May 23, 2011

## C++ at Google: Here Be Dragons

Google has one of the **largest monolithic C++ codebases** in the world. We have **thousands of engineers** working on millions of lines of C++ code every day.

<http://google-engtools.blogspot.com/2011/05/c-at-google-here-be-dragons.html>

# C++ Renaissance: Microsoft

---

**Pete Brown** is a XAML and Blinky lights guy at Microsoft who focuses on Windows XAML (WinRT), WPF, Silverlight, .NET Micro Framework and other "code on the client" and "code on a device" technologies. This is his personal blog.

Now, you can actually accomplish creating a basic add-in simply by storing a little hunk of script on the drive and adding the right registry keys. However, I specifically wanted to do this in C++. No, not because I hate myself, but because I'm starting to see a resurgence of interest in C++. You can create add-ins using .NET and Script, but both have significant limitations as well as performance concerns. If you want to write an add-in of any complexity, you'll almost certainly want to write it in C++. So, that's what I decided to do.

<http://10rem.net/blog/2011/02/22/creating-an-internet-explorer-add-in-toolbar-button-using-cplusplus-and-atl>

# C++ Renaissance: Facebook (2012)

---

Facebook consumes very little Wats per user. Actually fractions of Wat. ...  
Everything that is backend, high-performance is C++ because it must be responsive and well done. ...  
Facebook is literally running on C++. We also have JIT for PHP that is obviously written in C++ and generates machine code during runtime. ...  
Using the native language such as C++ means more features for the same power and less power for the given set of features. It is a very simple equation.

-- *Andrei Alexandrescu*, Research Scientist at Facebook, Lang.Next 2012

# C++ Renaissance: Facebook (Today)

---

It turns out **Facebook is already using JIT'ed C++ code in production** as their own **"efficient scripting framework"** for dealing with their HTTP request handling logic in their L7 reverse proxies. Their C++ scripts are compiled/linked/executed at run-time and built off the LLVM/Clang infrastructure.

Facebook engineers say their C++ scripting framework was **faster than another previously used scripting language by four times**. This isn't straight-up unmodified C++ sources they are JIT'ing but are relying upon pre-compiled headers and other changes to suit their workflow/requirements.

[https://www.phoronix.com/scan.php?page=news\\_item&px=Facebook-JIT-Cpp-Scripting](https://www.phoronix.com/scan.php?page=news_item&px=Facebook-JIT-Cpp-Scripting)

# C++ Renaissance: Bloomberg

---

Applications and Programming

## How Bloomberg is advancing C++ at scale



Bloomberg Professional August 23, 2016

<https://www.bloomberg.com/professional/blog/bloomberg-advancing-c-scale>

## WHAT IS C++?

# What is C++?

---

# What is C++?

---

- *C++ is no longer C with classes*

# What is C++?

---

- *C++ is no longer C with classes*
- C++ is a **general-purpose programming language**
  - *imperative, object-based, object-oriented, generic, and functional programming* features
  - developer is not tied to any specific programming paradigm by the language

# What is C++?

---

- *C++ is no longer C with classes*
- C++ is a **general-purpose programming language**
  - *imperative, object-based, object-oriented, generic, and functional programming* features
  - developer is not tied to any specific programming paradigm by the language
- **High-level abstractions at low cost**

# What is C++?

---

- *C++ is no longer C with classes*
- C++ is a **general-purpose programming language**
  - *imperative, object-based, object-oriented, generic, and functional programming* features
  - developer is not tied to any specific programming paradigm by the language
- **High-level abstractions at low cost**
- **Low-level access** when needed

# What is C++?

---

- *C++ is no longer C with classes*
- C++ is a **general-purpose programming language**
  - *imperative, object-based, object-oriented, generic, and functional programming* features
  - developer is not tied to any specific programming paradigm by the language
- **High-level abstractions at low cost**
- **Low-level access** when needed
- If used correctly, provides hard to beat **performance**

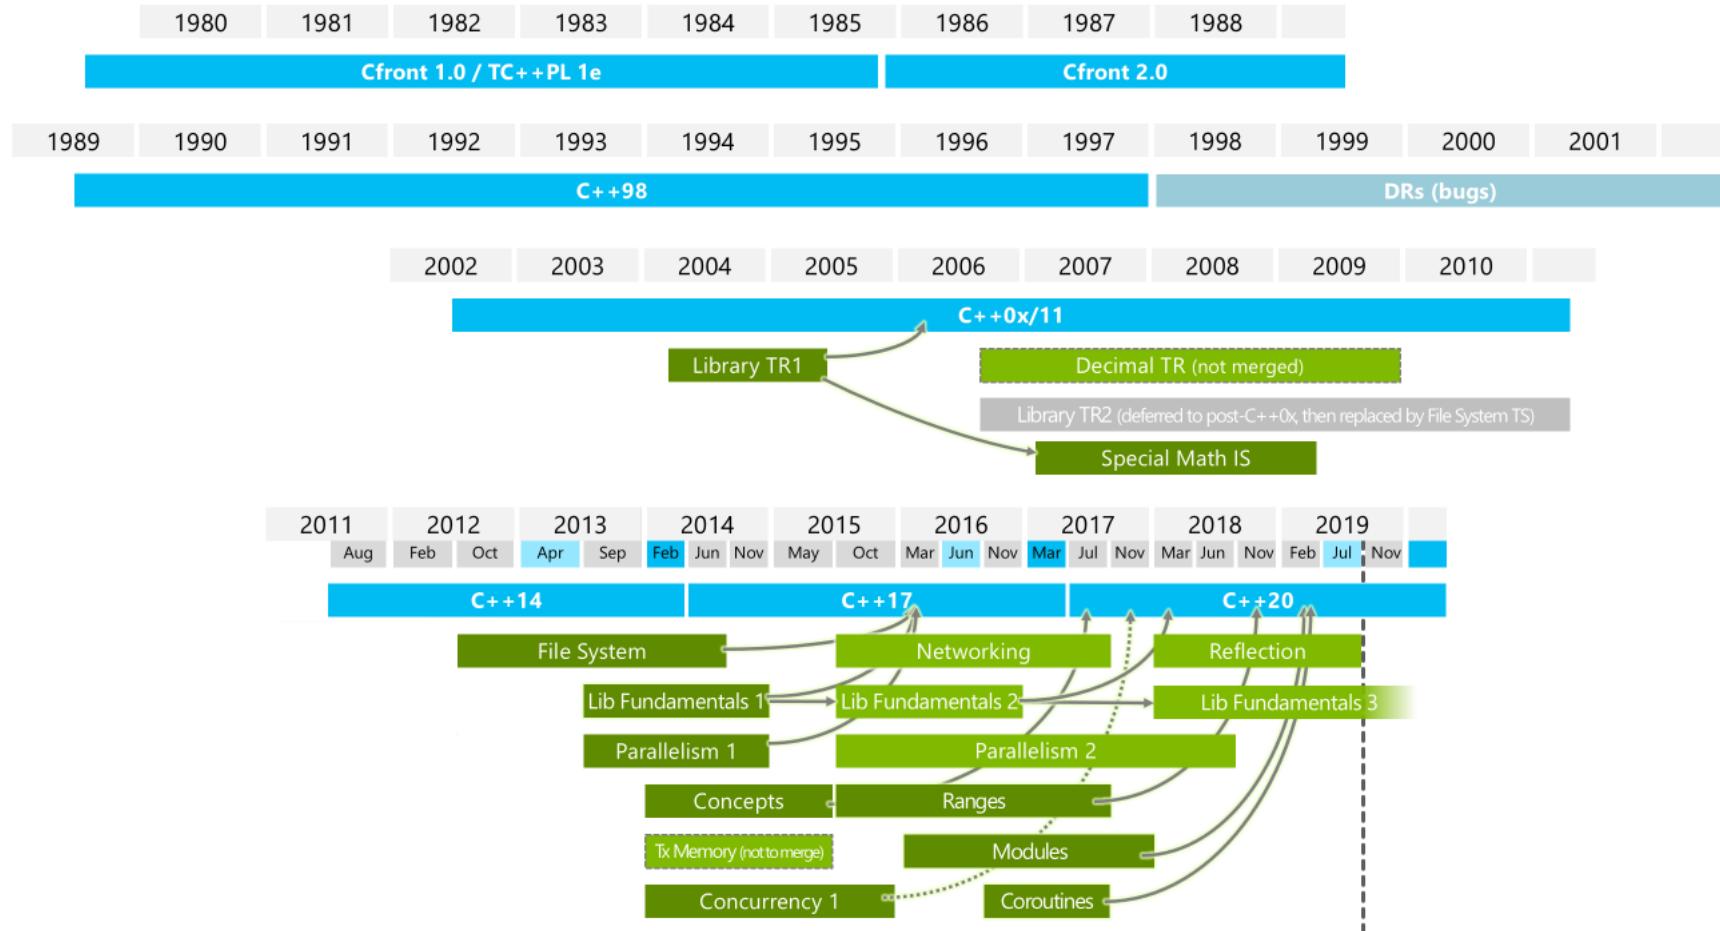
# What is C++?

---

- *C++ is no longer C with classes*
- C++ is a **general-purpose programming language**
  - *imperative, object-based, object-oriented, generic, and functional programming* features
  - developer is not tied to any specific programming paradigm by the language
- **High-level abstractions at low cost**
- **Low-level access** when needed
- If used correctly, provides hard to beat **performance**

C++ language evolves now faster than ever

# C++ Timeline



# ISO C++ Committee (WG21)

---

- ISO/IEC JTC1/SC22/WG21 (Joint Technical Committee 1/Subcommittee 22/Working Group 21)
- Formed in 1990-91
- Consists of accredited experts from member nations of ISO/IEC JTC1/SC22 who are interested in the C++ work
- Includes *13 member countries*: Bulgaria, Canada, Czech Republic, Finland, France, Germany, Netherlands, Poland, Russia, Spain, Switzerland, United Kingdom, and the United States
  - Bulgaria, Czech Republic, and Poland *joined during last year*

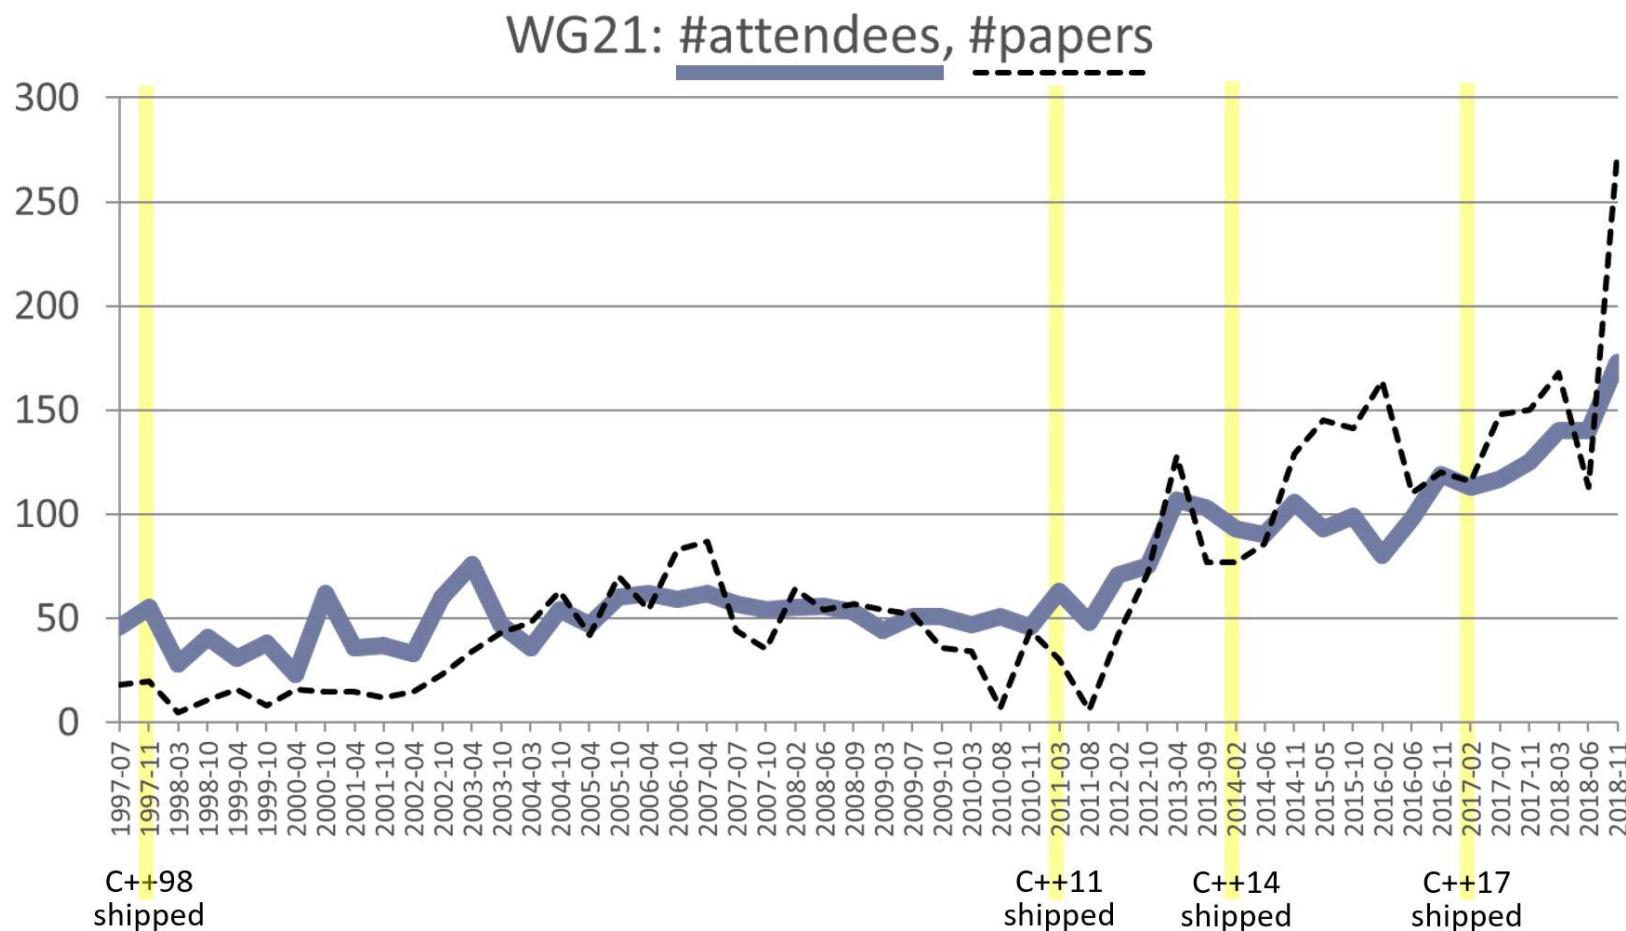
# ISO C++ Committee (WG21)

---

- ISO/IEC JTC1/SC22/WG21 (Joint Technical Committee 1/Subcommittee 22/Working Group 21)
- Formed in 1990-91
- Consists of accredited experts from member nations of ISO/IEC JTC1/SC22 who are interested in the C++ work
- Includes *13 member countries*: Bulgaria, Canada, Czech Republic, Finland, France, Germany, Netherlands, Poland, Russia, Spain, Switzerland, United Kingdom, and the United States
  - Bulgaria, Czech Republic, and Poland *joined during last year*

EPAM joined "The Committee" in February 2017!

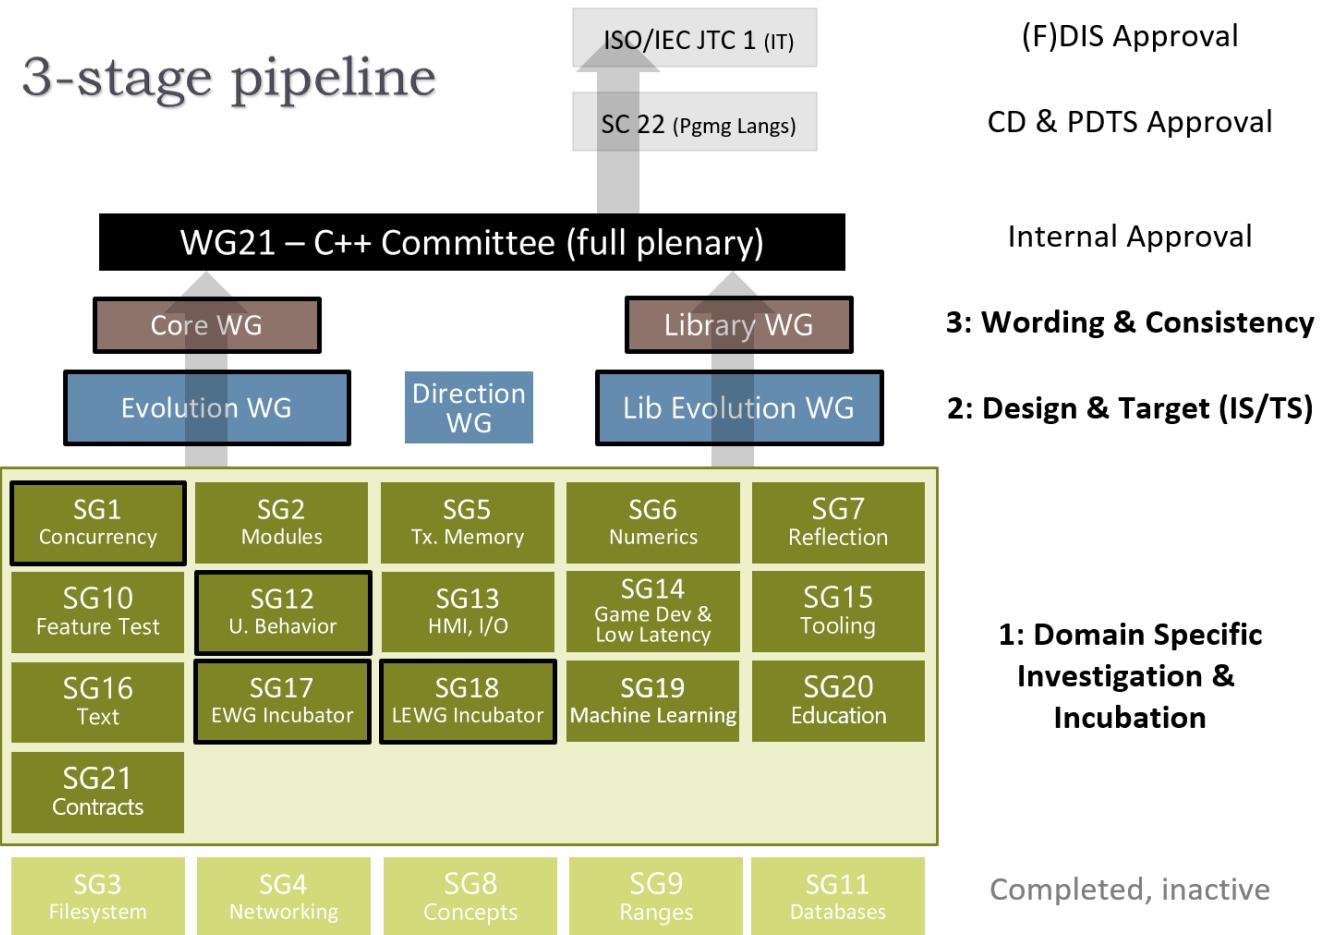
# C++ Momentum



# Not just yet another ISO committee

SC 22 SUBGROUP	NAME	RECENT FACE-TO-FACE ATTENDANCE	NOTES
WG 4	COBOL	15	
WG 5	Fortran	26	
WG 9	Ada	0	
WG 14	C	24	
WG 17	Prolog	3	
WG 21	C++	224	
WG 23	Programming Language Vulnerabilities	3	+ ~10 WG21 experts when meeting co-located with WG21
WG 24	Linux Standard Base	(none yet)	

# ISO C++ Committee structure



# Performance vs type safety and usability

---

# Performance vs type safety and usability

---

JAVA

```
class Price { /* ... */ };
```

```
Price buy = new Price(10);
Price threshold = new Price(11);
if(buy.less(threshold))
    execute();
buy.incr(1);
```

In most "modern" languages (i.e. Java, C#)

**user-defined types**

- have to be *allocated on the heap*
- do not have *value semantics*
- do not work with *built-in operators*

# Performance vs type safety and usability

JAVA

```
class Price { /* ... */ };
```

```
Price buy = new Price(10);
Price threshold = new Price(11);
if(buy.less(threshold))
    execute();
buy.incr(1);
```

C++

```
class price { /* ... */ };
```

```
price buy(10);
price threshold(11);
if(buy < threshold)
    execute();
buy += 1;
```

In most "modern" languages (i.e. Java, C#)

**user-defined types**

- have to be *allocated on the heap*
- do not have *value semantics*
- do not work with *built-in operators*

**Modern C++**

- *type safety without sacrificing performance*
- power and expressiveness of *built-in types*
- prefers *value semantics*
- *limits* dynamic allocations
- *discourages* dynamic polymorphism

# Because type matters

---

# Because type matters

---

C

```
void qsort(void* base, size_t el_count, size_t el_size,  
          int (*compare)(const void*, const void*));
```

# Because type matters

---

C

```
void qsort(void* base, size_t el_count, size_t el_size,  
          int (*compare)(const void*, const void*));
```

C++

```
template<class RandomIt, class Compare>  
constexpr void sort(RandomIt first, RandomIt last, Compare comp);
```

# Because type matters

---

C

```
void qsort(void* base, size_t el_count, size_t el_size,  
          int (*compare)(const void*, const void*));
```

C++

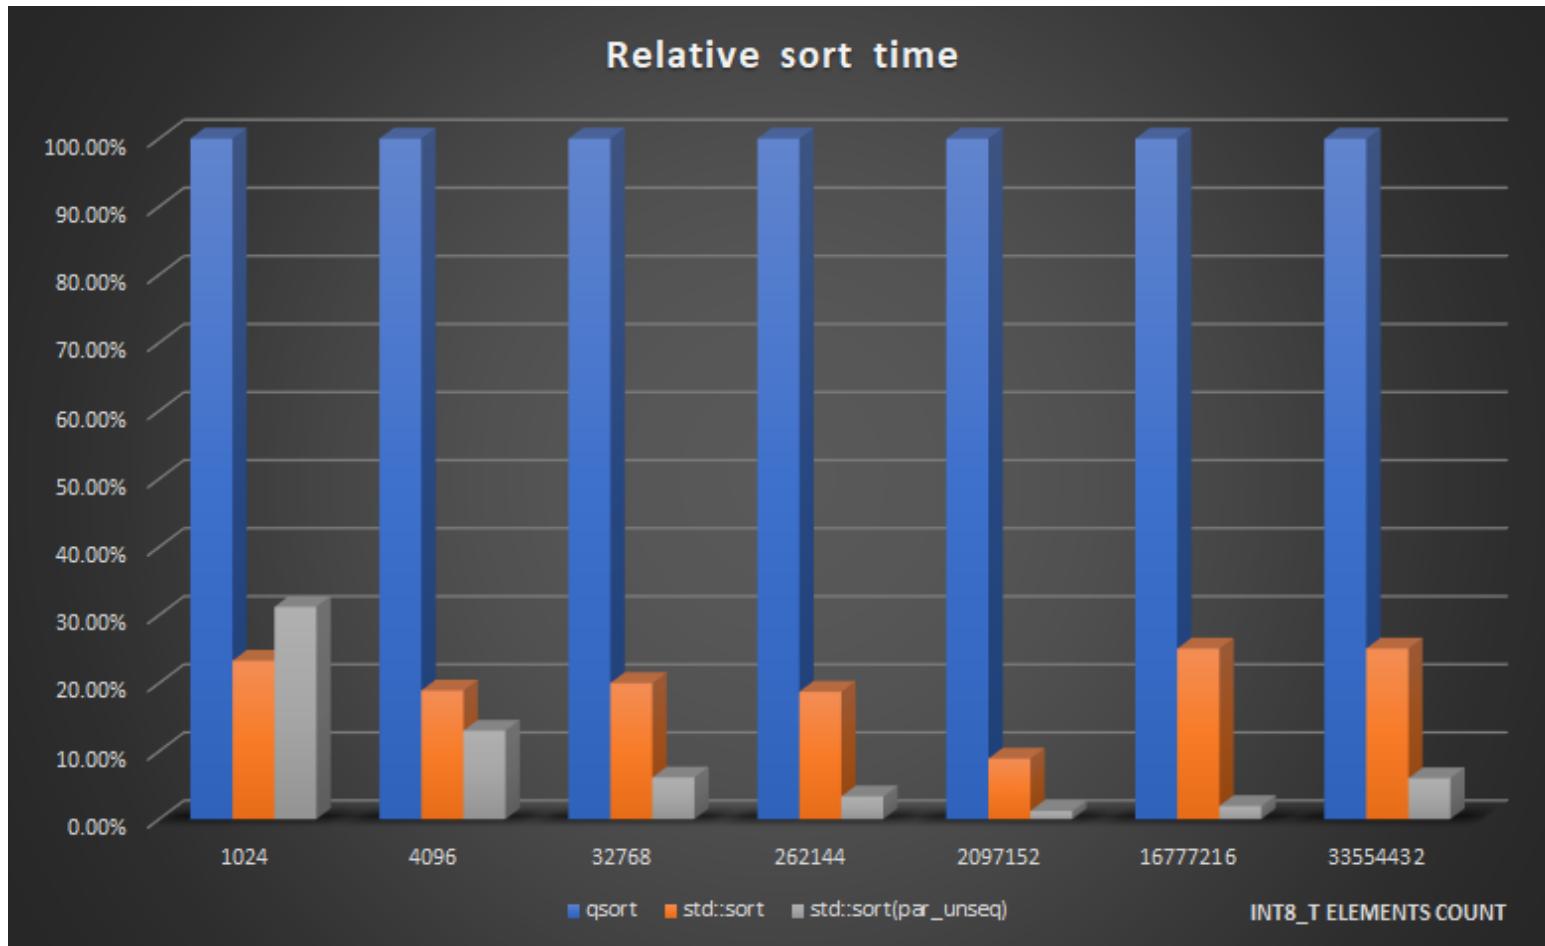
```
template<class RandomIt, class Compare>  
constexpr void sort(RandomIt first, RandomIt last, Compare comp);
```

```
template<class ExecutionPolicy, class RandomIt, class Compare>  
void sort(ExecutionPolicy&& policy, RandomIt first, RandomIt last, Compare comp);
```

# Because type matters



# Because type matters



# Unified approach to resource handling

---

# Unified approach to resource handling

---

## MEMORY

```
{  
    auto smart_ptr = std::make_unique<BigData>();  
    // use ptr  
}
```

# Unified approach to resource handling

---

## MEMORY

```
{  
    auto smart_ptr = std::make_unique<BigData>();  
    // use ptr  
}
```

## THREAD SYNCHRONIZATION

```
{  
    auto lock = std::scoped_lock(mtx);  
    // critical section  
}
```

# Unified approach to resource handling

## MEMORY

```
{  
    auto smart_ptr = std::make_unique<BigData>();  
    // use ptr  
}
```

## FILE HANDLES

```
{  
    auto file = fhandle(fopen("demo.txt", "r"));  
    // use the file  
}
```

## THREAD SYNCHRONIZATION

```
{  
    auto lock = std::scoped_lock(mutex);  
    // critical section  
}
```

# Unified approach to resource handling

## MEMORY

```
{  
    auto smart_ptr = std::make_unique<BigData>();  
    // use ptr  
}
```

## FILE HANDLES

```
{  
    auto file = fhandle(fopen("demo.txt", "r"));  
    // use the file  
}
```

## THREAD SYNCHRONIZATION

```
{  
    auto lock = std::scoped_lock(mutex);  
    // critical section  
}
```

## NETWORK SOCKET

```
{  
    auto socket = connect("www.epam.com");  
    // use the socket  
}
```

# Unified approach to resource handling

## MEMORY

```
{  
    auto smart_ptr = std::make_unique<BigData>();  
    // use ptr  
}
```

## FILE HANDLES

```
{  
    auto file = fhandle(fopen("demo.txt", "r"));  
    // use the file  
}
```

## THREAD SYNCHRONIZATION

```
{  
    auto lock = std::scoped_lock(mutex);  
    // critical section  
}
```

## NETWORK SOCKET

```
{  
    auto socket = connect("www.epam.com");  
    // use the socket  
}
```

No need for a garbage collector as memory is not the only resource to handle

# Compile time safety without sacrificing performance

---

# Compile time safety without sacrificing performance

```
constexpr double avg_speed(double length,
                           double time)
{
    return length / time;
}
```

```
/// l - distance in kilometers
/// t - time in hours
void foo(double l, double t)
{
    auto speed = avg_speed(l, t);
    int mps = int(speed * 1000 / 3600);

    printf("Average speed in mps is: %u\n",
           mps);
}
```

# Compile time safety without sacrificing performance

```
constexpr double avg_speed(double length,  
                           double time)  
{  
    return length / time;  
}
```

```
/// l - distance in kilometers  
/// t - time in hours  
void foo(double l, double t)  
{  
    auto speed = avg_speed(l, t);  
    int mps = int(speed * 1000 / 3600);  
  
    printf("Average speed in mps is: %u\n",  
          mps);  
}
```

```
constexpr Velocity auto avg_speed(Length auto d,  
                                  Time auto t)  
{  
    return d / t;  
}
```

```
void foo(Length auto l, Time auto t)  
{  
    auto speed = avg_speed(l, t);  
    auto mps = quantity_cast<metre_per_second,  
                           int>(speed);  
    printf("Average speed in mps is: %u\n",  
          mps.count());  
}
```

# Compile time safety without sacrificing performance

```
constexpr double avg_speed(double length,  
                           double time)  
{  
    return length / time;  
}
```

```
/// l - distance in kilometers  
/// t - time in hours  
void foo(double l, double t)  
{  
    auto speed = avg_speed(l, t);  
    int mps = int(speed * 1000 / 3600);  
  
    printf("Average speed in mps is: %u\n",  
          mps);  
}
```

```
constexpr Velocity auto avg_speed(Length auto d,  
                                  Time auto t)  
{  
    return d / t;  
}
```

```
void foo(Length auto l, Time auto t)  
{  
    auto speed = avg_speed(l, t);  
    auto mps = quantity_cast<metre_per_second,  
                           int>(speed);  
    printf("Average speed in mps is: %u\n",  
          mps.count());  
}
```

How much do we pay for such advanced and safe abstractions?

# Compile time safety without sacrificing performance

```
1 #include <cstdio>
2
3
4
5
6 constexpr double avg_speed(double length, double time)
7 {
8     return length / time;
9 }
10
11 void foo(double l, double t)
12 {
13     auto speed = avg_speed(l, t);
14     printf("Average speed in mps is: %u\n",
15           int(speed * 1000 / 3600));
16 }
17
18 int main()
19 {
20     foo(144., 2.);
21 }
22
```

The screenshot shows the Compiler Explorer interface with the following details:

- Compiler:** x86-64 gcc 9.2
- Flags:** -std=c++2a -fconcepts -O2 -DNDEBUG
- Output:**
  - Assembly code:

```
1 .LC2:
2     .string "Average speed in mps is: %u\n"
3 foo(double, double):
4     divsd    xmm0, xmm1
5     mov      edi, OFFSET FLAT:.LC2
6     xor      eax, eax
7     mulsd   xmm0, QWORD PTR .LC0[rip]
8     divsd   xmm0, QWORD PTR .LC1[rip]
9     cvttsd2si    esi, xmm0
10    jmp     printf
11 main:
12    sub     rsp, 8
13    mov     esi, 20
14    mov     edi, OFFSET FLAT:.LC2
```
  - Output window:

```
ASM generation compiler returned: 0
Execution build compiler returned: 0
Program returned: 0
Average speed in mps is: 20
```
- Bottom right corner:** Edit on Compiler Explorer

# Compile time safety without sacrificing performance

```
1 #include <units/dimensions/velocity.h>
2
3 using namespace units;
4
5 template<Length L, Time T>
6 constexpr Velocity avg_speed(L length, T time)
7 {
8     return length / time;
9 }
10
11 void foo(double l, double t)
12 {
13     auto speed = avg_speed(quantity<kilometre>(1), quantity<hour>(t))
14     printf("Average speed in mps is: %u\n",
15           quantity_cast<metre_per_second, int>(speed).count());
16 }
17
18 int main()
19 {
20     foo(144., 2.);
21 }
22
```

The screenshot shows the Compiler Explorer interface with the following configuration:

- Compiler: x86-64 gcc 9.2
- Flags: -std=c++2a -fconcepts -O2 -DNDEBUG
- Output tabs: A (selected), 11010, ./a.out, .LX0:, lib.f., .text, //, \s+, Intel, Demangle, Libraries
- Tool buttons: + Add new..., Add tool...

The assembly code generated by the compiler is as follows:

```
1 .LC2:
2     .string "Average speed in mps is: %u\n"
3     foo(double, double):
4         divsd    xmm0, xmm1
5         mov      edi, OFFSET FLAT:.LC2
6         xor      eax, eax
7         mulsd    xmm0, QWORD PTR .LC0[rip]
8         divsd    xmm0, QWORD PTR .LC1[rip]
9         cvttsd2si    esi, xmm0
10        jmp     printf
11     main:
12         sub     rsp, 8
13         mov     esi, 20
14         mov     edi, OFFSET FLAT:.LC2
```

The output window shows the following results:

```
ASM generation compiler returned: 0
Execution build compiler returned: 0
Program returned: 0
Average speed in mps is: 20
```

Bottom right corner: Edit on Compiler Explorer

# Compile time safety without sacrificing performance

```
1 #include <units/dimensions/velocity.h>
2
3 using namespace units;
4
5 template<Length L, Time T>
6 constexpr Velocity avg_speed(L length, T time)
7 {
8     return length / time;
9 }
10
11 void foo(Length l, Time t)
12 {
13     auto speed = avg_speed(l, t);
14     printf("Average speed in mps is: %u\n",
15           quantity_cast<metre_per_second, int>(speed).count());
16 }
17
18 int main()
19 {
20     using namespace units::literals;
21     foo(144.km, 2.h);
22 }
```

The screenshot shows the Compiler Explorer interface with the following details:

- Compiler:** x86-64 gcc 9.2
- Flags:** -std=c++2a -fconcepts -O2 -DNDEBUG
- Output:** Assembly code generated by the compiler.
- Assembly Output:**

```
1 .LC0:
2     .string "Average speed in mps is: %u\n"
3 main:
4     sub    rsp, 8
5     mov    esi, 20
6     mov    edi, OFFSET FLAT:.LC0
7     xor    eax, eax
8     call   printf
9     xor    eax, eax
10    add   rsp, 8
11    ret
```
- Execution Results:**
  - ASM generation compiler returned: 0
  - Execution build compiler returned: 0
  - Program returned: 0
  - Average speed in mps is: 20
- Bottom Right:** Edit on Compiler Explorer

# Compile time safety without sacrificing performance

The screenshot shows a developer environment with two main panes. The left pane displays the C++ source code:

```
1 #include <cstdio>
2
3
4
5
6
7
8
9
10
11
12
13
14
15
16
17
18 int main()
19 {
20     printf("Average speed in mps is: %u\n", 20);
21 }
22
```

The right pane shows the assembly output and the terminal output. The assembly code is:

```
1 .LC0:
2     .string "Average speed in mps is: %u\n"
3 main:
4     sub    rsp, 8
5     mov    esi, 20
6     mov    edi, OFFSET FLAT:.LC0
7     xor    eax, eax
8     call   printf
9     xor    eax, eax
10    add   rsp, 8
11    ret
```

The terminal output is:

```
ASM generation compiler returned: 0
Execution build compiler returned: 0
Program returned: 0
Average speed in mps is: 20
```

At the bottom right of the terminal window, there is a link: [Edit on Compiler Explorer](#).

# Omnipresence of compile-time usage

---

# Omnipresence of compile-time usage

---

## DATE HANDLING

```
constexpr auto d = 2019_y/September/21;  
static_assert(weekday{d} == Saturday);
```

# Omnipresence of compile-time usage

---

## DATE HANDLING

```
constexpr auto d = 2019_y/September/21;  
static_assert(weekday{d} == Saturday);
```

```
std::cout << "'SEC 2019' starts" <<  
    " on " << weekday{d} <<  
    " at " << d + 9h << "\n";
```

# Omnipresence of compile-time usage

---

## DATE HANDLING

```
constexpr auto d = 2019_y/September/21;  
static_assert(weekday{d} == Saturday);
```

```
std::cout << "'SEC 2019' starts" <<  
    " on " << weekday{d} <<  
    " at " << d + 9h << "\n";
```

```
static_assert(d == September/21/2019);  
static_assert(d == 21_d/September/2019);  
static_assert(d == Saturday[3]/September/2019);
```

# Omnipresence of compile-time usage

## DATE HANDLING

```
constexpr auto d = 2019_y/September/21;  
static_assert(weekday{d} == Saturday);
```

```
std::cout << "'SEC 2019' starts" <<  
    " on " << weekday{d} <<  
    " at " << d + 9h << "\n";
```

```
static_assert(d == September/21/2019);  
static_assert(d == 21_d/September/2019);  
static_assert(d == Saturday[3]/September/2019);
```

## UNITS AND DIMENSIONAL ANALYSIS

```
// simple numeric operations  
static_assert(10km / 2 == 5km);  
  
// unit conversions  
static_assert(1h == 3600s);  
static_assert(1km + 1m == 1001m);  
  
// dimension conversions  
static_assert(1km / 1s == 1000mps);  
static_assert(2kmph * 2h == 4km);  
static_assert(2km / 2kmph == 1h);  
static_assert(1000 / 1s == 1kHz);  
static_assert(10km / 5km == 2);
```

# Omnipresence of compile-time usage

## DATE HANDLING

```
constexpr auto d = 2019_y/September/21;  
static_assert(weekday{d} == Saturday);
```

```
std::cout << "'SEC 2019' starts" <<  
    " on " << weekday{d} <<  
    " at " << d + 9h << "\n";
```

```
static_assert(d == September/21/2019);  
static_assert(d == 21_d/September/2019);  
static_assert(d == Saturday[3]/September/2019);
```

## UNITS AND DIMENSIONAL ANALYSIS

```
// simple numeric operations  
static_assert(10km / 2 == 5km);
```

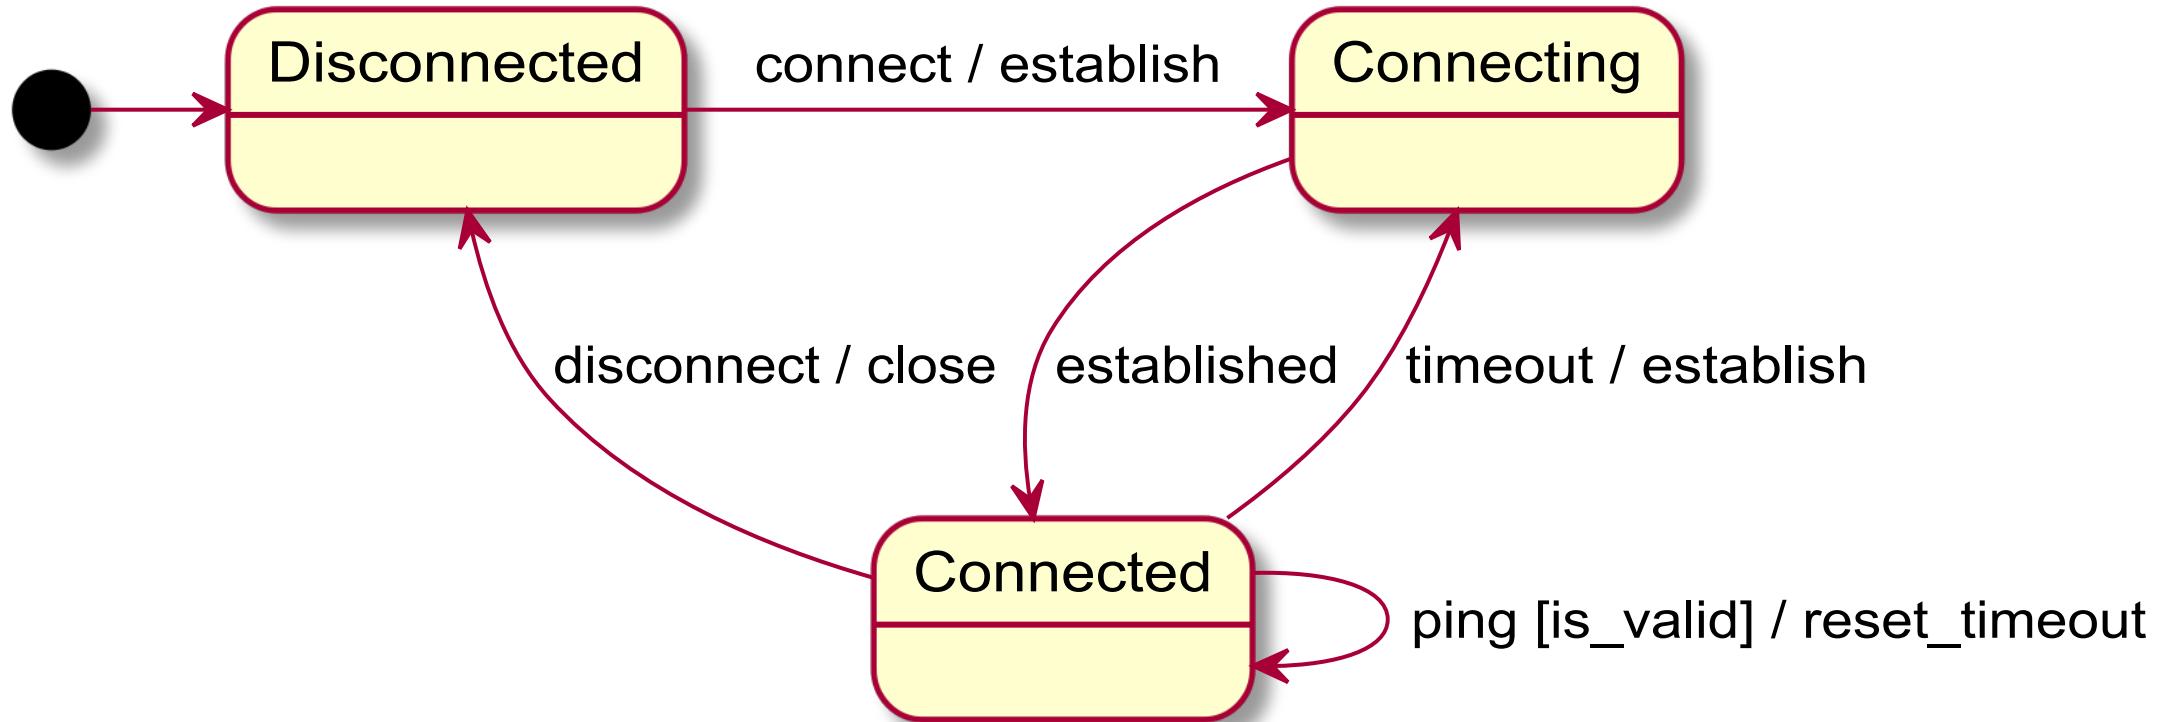
```
// unit conversions  
static_assert(1h == 3600s);  
static_assert(1km + 1m == 1001m);
```

```
// dimension conversions  
static_assert(1km / 1s == 1000mps);  
static_assert(2kmph * 2h == 4km);  
static_assert(2km / 2kmph == 1h);  
static_assert(1000 / 1s == 1kHz);  
static_assert(10km / 5km == 2);
```

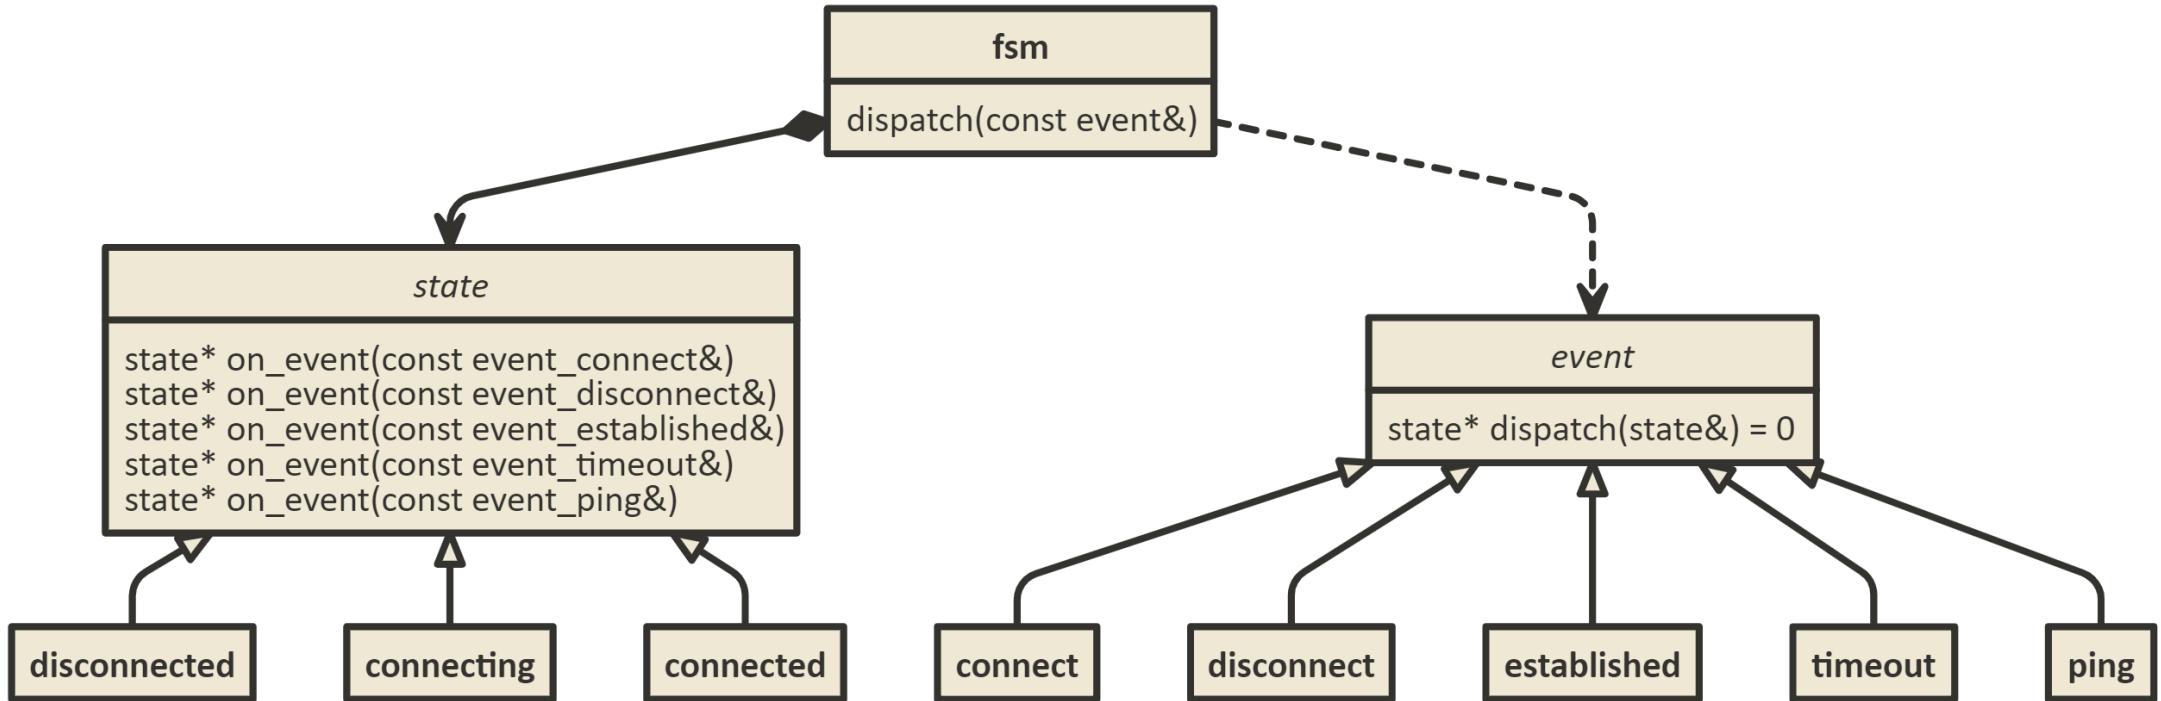
The fastest programs are those that do nothing

# Finite State Machine (FSM)

---



# FSM: A classical approach

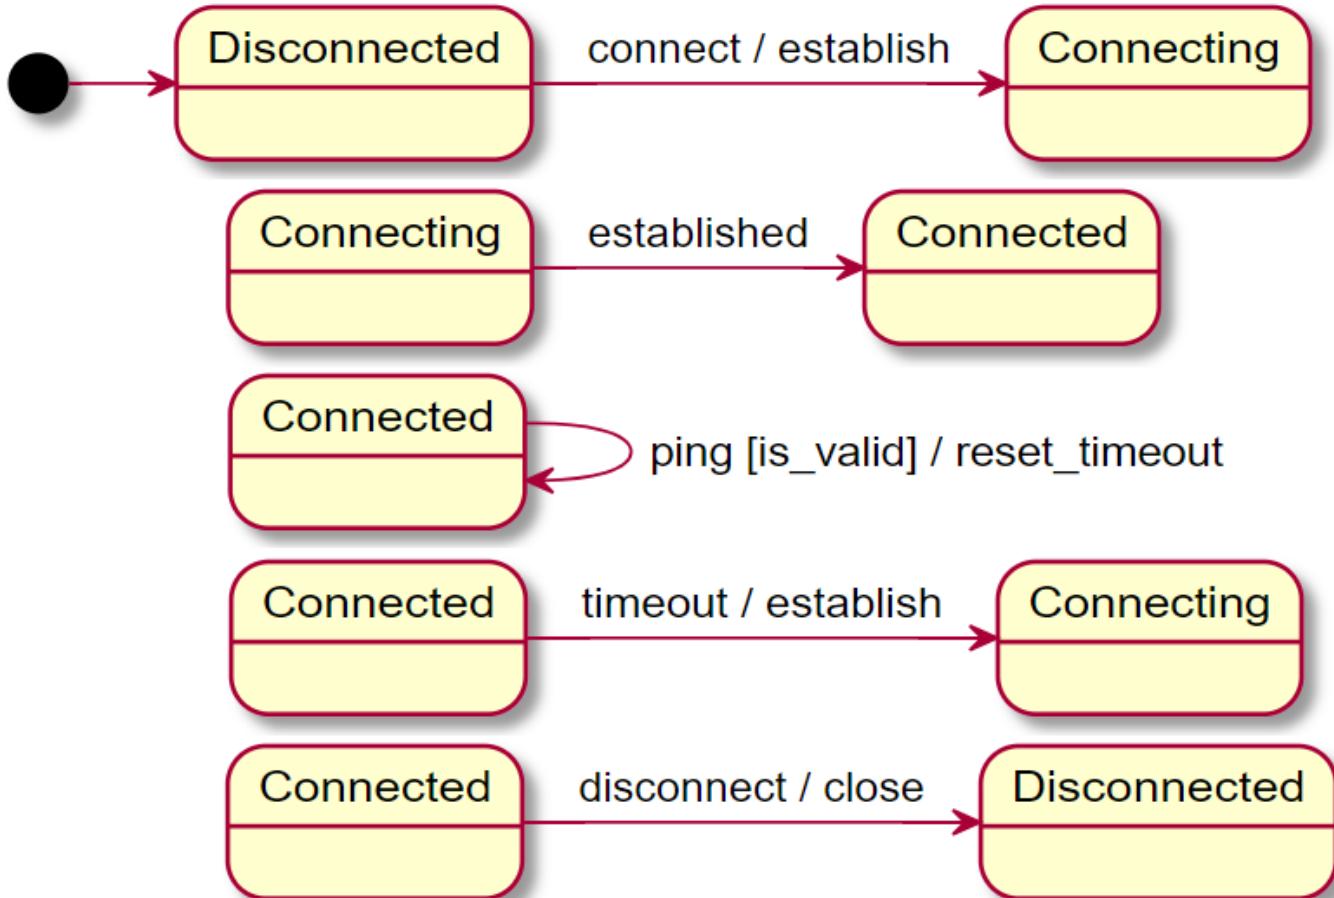


# FSM: A naive approach

```
class fsm {
    enum class state : char { disconnected, connecting, connected } state_;
public:
    constexpr void on_event(const connect&)
    {
        if(state_ == state::disconnected) {
            establish();
            state_ = state::connecting;
        }
    }

    constexpr void on_event(const disconnect&) { /* ... */ }
    constexpr void on_event(const established&) { /* ... */ }
    constexpr void on_event(const ping& event) { /* ... */ }
    constexpr void on_event(const timeout&) { /* ... */ }
};
```

# FSM: Alternative syntax



# FSM: Domain Specific Language

```
struct connection {
    auto operator()() const
    {
        using namespace boost::sml;
        return make_transition_table(
            * "Disconnected"_s + event<connect> / establish      = "Connecting"_s,
            "Connecting"_s   + event<established>                 = "Connected"_s,
            "Connected"_s    + event<ping> [ is_valid ] / reset_timeout,
            "Connected"_s    + event<timeout> / establish       = "Connecting"_s,
            "Connected"_s    + event<disconnect> / close         = "Disconnected"_s
        );
    }
};
```

# FSM: Domain Specific Language

```
struct connection {
    auto operator()() const
    {
        using namespace boost::sml;
        return make_transition_table(
            * "Disconnected"_s + event<connect> / establish      = "Connecting"_s,
            "Connecting"_s   + event<established>                 = "Connected"_s,
            "Connected"_s    + event<ping> [ is_valid ] / reset_timeout,
            "Connected"_s    + event<timeout> / establish       = "Connecting"_s,
            "Connected"_s    + event<disconnect> / close         = "Disconnected"_s
        );
    }
};
```

Program that says WHAT rather than HOW is more likely to be correct

# FSM: A classical approach

```
17 // state base class
18 struct state {
19     virtual ~state() = default;
20     virtual std::unique_ptr<state> on_event(const connect&) { return nullptr; }
21     virtual std::unique_ptr<state> on_event(const ping&) { return nullptr; }
22     virtual std::unique_ptr<state> on_event(const established &) { return nullptr; }
23     virtual std::unique_ptr<state> on_event(const timeout&) { return nullptr; }
24     virtual std::unique_ptr<state> on_event(const disconnect&) { return nullptr; }
25 };
26
27 // events
28 struct event {
29     virtual ~event() = default;
30     virtual std::unique_ptr<state> dispatch(state&) const = 0;
31 };
32
33 struct connect final : event {
34     std::unique_ptr<state> dispatch(state& s) const override { return s.on_event(*this); }
35 };
36
37 struct established final : event {
38     std::unique_ptr<state> dispatch(state& s) const override { return s.on_event(*this); }
39 };
40
41 struct ping final : event {
42     std::unique_ptr<state> dispatch(state& s) const override { return s.on_event(*this); }
43 };
44
45 struct disconnect final : event {
46     std::unique_ptr<state> dispatch(state& s) const override { return s.on_event(*this); }
47 };
48
49 struct timeout final : event {
50     std::unique_ptr<state> dispatch(state& s) const override { return s.on_event(*this); }
```

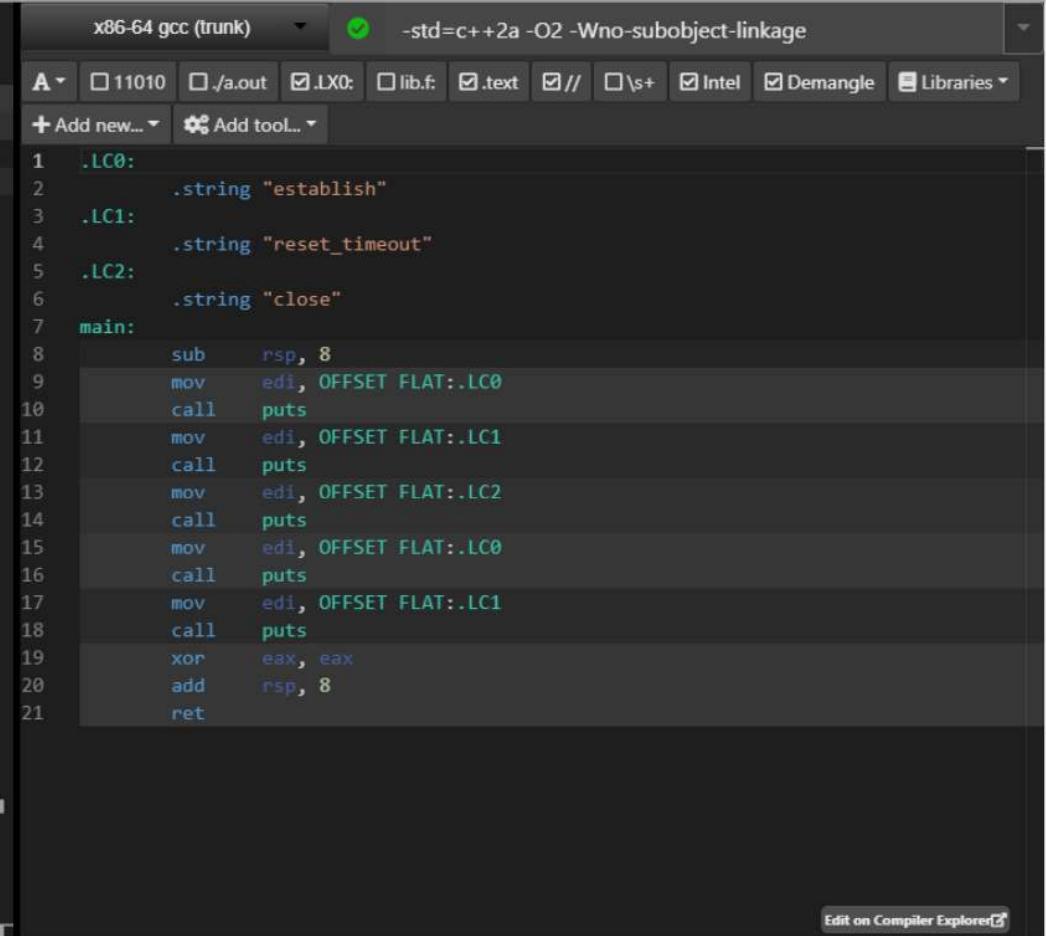
The screenshot shows the Compiler Explorer interface for an x86-64 gcc (trunk) build with flags `-std=c++2a -O2 -Wno-subobject-linkage`. The assembly output is displayed in two columns: address (322, 323, etc.) and assembly code.

Key assembly output:

- Address 322: .quad 0
- Address 323: .quad typeinfo for connected
- Address 324: .quad connected::~connected() [complete object destructor]
- Address 325: .quad connected::~connected() [deleting destructor]
- Address 326: state::on\_event(connect const&)
- Address 327: connected::on\_event(ping const&)
- Address 328: state::on\_event(established const&)
- Address 329: connected::on\_event(timeout const&)
- Address 330: connected::on\_event(disconnect const&)
- Address 331: vtable for disconnected:
  - Address 332: .quad 0
  - Address 333: .quad typeinfo for disconnected
  - Address 334: .quad disconnected::~disconnected() [complete object destructor]
  - Address 335: .quad disconnected::~disconnected() [deleting destructor]
  - Address 336: disconnected::on\_event(connect const&)
  - Address 337: state::on\_event(ping const&)
  - Address 338: state::on\_event(established const&)
  - Address 339: state::on\_event(timeout const&)
  - Address 340: state::on\_event(disconnect const&)
- Address 341: vtable for connecting:
  - Address 342: .quad 0
  - Address 343: .quad typeinfo for connecting
  - Address 344: .quad connecting::~connecting() [complete object destructor]
  - Address 345: .quad connecting::~connecting() [deleting destructor]
  - Address 346: state::on\_event(connect const&)
  - Address 347: state::on\_event(ping const&)
  - Address 348: connecting::on\_event(established const&)
  - Address 349: state::on\_event(timeout const&)
  - Address 350: state::on\_event(disconnect const&)

# FSM: A naive approach

```
1 #include <cstdio>
2
3 // actions
4 constexpr auto establish = []{ std::puts("establish"); };
5 constexpr auto close = []{ std::puts("close"); };
6 constexpr auto is_valid = [](auto const&) { return true; };
7 constexpr auto reset_timeout = []{ std::puts("reset_timeout"); };
8
9 // events
10 struct connect {};
11 struct established {};
12 struct ping {};
13 struct disconnect {};
14 struct timeout {};
15
16 // fsm
17 class fsm {
18     enum class state : char { disconnected, connecting, connected } state_;
19 public:
20     constexpr void on_event(const connect&)
21     {
22         if(state_ == state::disconnected) {
23             establish();
24             state_ = state::connecting;
25         }
26     }
27
28     constexpr void on_event(const disconnect&)
29     {
30         if(state_ == state::connecting || state_ == state::connected) {
31             close();
32             state_ = state::disconnected;
33         }
34     }
35 }
```



The screenshot shows the assembly output for the provided C++ code. The assembly code is generated for an x86-64 architecture using gcc (trunk) with flags -std=c++2a -O2 -Wno-subobject-linkage. The assembly code consists of several labels and instructions:

- Labels: .LC0, .LC1, .LC2, main.
- Instructions:
  - .LC0: .string "establish"
  - .LC1: .string "reset\_timeout"
  - .LC2: .string "close"
  - main: sub rsp, 8
  - mov edi, OFFSET FLAT:.LC0
  - call puts
  - mov edi, OFFSET FLAT:.LC1
  - call puts
  - mov edi, OFFSET FLAT:.LC2
  - call puts
  - mov edi, OFFSET FLAT:.LC0
  - call puts
  - xor eax, eax
  - add rsp, 8
  - ret

The assembly code corresponds to the logic in the C++ code: it prints "establish", "reset\_timeout", and "close" in sequence. The compiler has optimized the code by using global string literals (.LC0, .LC1, .LC2) and calling the standard library's puts function for each string.

# FSM: Domain Specific Language

```
1 #include <https://raw.githubusercontent.com/boost-experimental/sml/master/include/boost/sml.hpp>
2 #include <iostream>
3
4 // actions
5 constexpr auto establish = []{ std::puts("establish"); };
6 constexpr auto close = []{ std::puts("close"); };
7 constexpr auto is_valid = [](auto const&){ return true; };
8 constexpr auto reset_timeout = []{ std::puts("reset_timeout"); };
9
10 // states
11 struct connect {};
12 struct established {};
13 struct ping {};
14 struct disconnect {};
15 struct timeout {};
16
17 // fsm
18 struct connection {
19     auto operator()() const
20     {
21         using namespace boost::sml;
22         return make_transition_table(
23             * "Disconnected"_s + event<connect> / establish
24             = "Connecting"_s,
25             "Connecting"_s + event<established>
26             = "Connected"_s,
27             "Connected"_s + event<ping> [ is_valid ] / reset_timeout,
28             = "Connecting"_s,
29             "Connected"_s + event<disconnect> / close
30             = "Disconnected"_s
31         );
32     }
33     int main()
34     {
35         boost::sml::sm<connection> connection{};
```

x86-64 gcc (trunk) -std=c++2a -O2 -Wno-subobject-linkage

A 11010 ./a.out LX0 libf. text // \s+ Intel Demangle Libraries

+ Add new... Add tool...

```
1 .LC0:
2     .string "establish"
3 .LC1:
4     .string "reset_timeout"
5 .LC2:
6     .string "close"
7 main:
8     sub    rsp, 8
9     mov    edi, OFFSET FLAT:.LC0
10    call   puts
11    mov    edi, OFFSET FLAT:.LC1
12    call   puts
13    mov    edi, OFFSET FLAT:.LC2
14    call   puts
15    mov    edi, OFFSET FLAT:.LC0
16    call   puts
17    mov    edi, OFFSET FLAT:.LC1
18    call   puts
19    xor    eax, eax
20    add    rsp, 8
21    ret
```

Edit on Compiler Explorer

# FSM: Summary

---

	CLASSICAL	NAIVE	DSL
HEAP USAGE	Lots	None	None
DYNAMIC POLYMORPHISM	Yes	No	No
INLINEING	No	Yes	Yes
CPU CACHE FRIENDLY	No	Yes	Yes
ERROR PRONE	Yes	Yes	No
REUSE	Moderate	Hard	Easy
ASSEMBLY LINES	350	21	21
MEMORY FOOTPRINT	High	1B	1B
PERFORMANCE	SLOW	FAST	FAST

# The best language for Data Oriented Design (DOD)

---

# The best language for Data Oriented Design (DOD)

---

## 1 Keep data that is not used together apart, on separate cache lines

- avoids the invisible convoying of false sharing (ping-pong)

# The best language for Data Oriented Design (DOD)

---

1 **Keep data that is not used together apart**, on separate cache lines

- avoids the invisible convoying of false sharing (ping-pong)

2 **Keep data that is frequently used together close together**

- if a thread that uses A frequently also needs B, try to put them in one cache line

# The best language for Data Oriented Design (DOD)

---

1 **Keep data that is not used together apart**, on separate cache lines

- avoids the invisible convoying of false sharing (ping-pong)

2 **Keep data that is frequently used together close together**

- if a thread that uses A frequently also needs B, try to put them in one cache line

3 **Keep “hot” and “cold” data that is not used with the same frequency apart**

- fit “hot” data in the fewest possible cache lines and memory pages
- reduces the cache footprint and cache misses, the memory footprint, and virtual memory paging

# The best language for Data Oriented Design (DOD)

---

1 Keep data that is not used together apart, on separate cache lines

- avoids the invisible convoying of false sharing (ping-pong)

2 Keep data that is frequently used together close together

- if a thread that uses A frequently also needs B, try to put them in one cache line

3 Keep “hot” and “cold” data that is not used with the same frequency apart

- fit “hot” data in the fewest possible cache lines and memory pages
- reduces the cache footprint and cache misses, the memory footprint, and virtual memory paging

CPU instructions to execute is also data sitting in L1 cache

Modern C++ is not the language we used to learn 20 years ago!

Modern C++ is not the language we used to learn 20 years ago!

- **Move Semantics** to efficiently handle resources and improve performance

Modern C++ is not the language we used to learn 20 years ago!

- **Move Semantics** to efficiently handle resources and improve performance
- **constexpr** to move execution of the code from run-time to compile-time

Modern C++ is not the language we used to learn 20 years ago!

- **Move Semantics** to efficiently handle resources and improve performance
- **constexpr** to move execution of the code from run-time to compile-time
- **Deterministic** destruction

Modern C++ is not the language we used to learn 20 years ago!

- **Move Semantics** to efficiently handle resources and improve performance
- **constexpr** to move execution of the code from run-time to compile-time
- **Deterministic** destruction
- Better **resource handling**

Modern C++ is not the language we used to learn 20 years ago!

- **Move Semantics** to efficiently handle resources and improve performance
- **constexpr** to move execution of the code from run-time to compile-time
- **Deterministic** destruction
- Better **resource handling**
- **Value semantics** to improve performance and development experience

Modern C++ is not the language we used to learn 20 years ago!

- **Move Semantics** to efficiently handle resources and improve performance
- **constexpr** to move execution of the code from run-time to compile-time
- **Deterministic** destruction
- Better **resource handling**
- **Value semantics** to improve performance and development experience
- Full control over **concurrency and parallelism** (task-based, parallel, vectorized, lock-free, ...)

Modern C++ is not the language we used to learn 20 years ago!

- Move Semantics to efficiently handle resources and improve performance
- `constexpr` to move execution of the code from run-time to compile-time
- Deterministic destruction
- Better resource handling
- Value semantics to improve performance and development experience
- Full control over concurrency and parallelism (task-based, parallel, vectorized, lock-free, ...)
- Many productivity-related improvements (i.e. lambdas, range-for loops, `auto`, variadic templates, ...)

Modern C++ is not the language we used to learn 20 years ago!

- Move Semantics to efficiently handle resources and improve performance
- `constexpr` to move execution of the code from run-time to compile-time
- Deterministic destruction
- Better resource handling
- Value semantics to improve performance and development experience
- Full control over concurrency and parallelism (task-based, parallel, vectorized, lock-free, ...)
- Many productivity-related improvements (i.e. lambdas, range-for loops, `auto`, variadic templates, ...)
- Lots of new toys in the C++ Standard Library

# C++20 is going to be a game-changer

---

# C++20 is going to be a game-changer

---

- Concepts to provide even better compile-time safety

# C++20 is going to be a game-changer

---

- **Concepts** to provide even better compile-time safety
- **Coroutines** to enhance asynchronous operations

# C++20 is going to be a game-changer

---

- **Concepts** to provide even better compile-time safety
- **Coroutines** to enhance asynchronous operations
- **Ranges** to build better algorithms with higher level abstractions

# C++20 is going to be a game-changer

---

- **Concepts** to provide even better compile-time safety
- **Coroutines** to enhance asynchronous operations
- **Ranges** to build better algorithms with higher level abstractions
- **Modules** to fix C build system and improve large-scale architecture

# C++20 is going to be a game-changer

---

- **Concepts** to provide even better compile-time safety
- **Coroutines** to enhance asynchronous operations
- **Ranges** to build better algorithms with higher level abstractions
- **Modules** to fix C build system and improve large-scale architecture
- **<=>** to simplify implementation of value types

# C++20 is going to be a game-changer

---

- **Concepts** to provide even better compile-time safety
- **Coroutines** to enhance asynchronous operations
- **Ranges** to build better algorithms with higher level abstractions
- **Modules** to fix C build system and improve large-scale architecture
- **<=>** to simplify implementation of value types
- **New tools** in the C++ Standard Library

# C++20 is going to be a game-changer

---

- **Concepts** to provide even better compile-time safety
- **Coroutines** to enhance asynchronous operations
- **Ranges** to build better algorithms with higher level abstractions
- **Modules** to fix C build system and improve large-scale architecture
- **<=>** to simplify implementation of value types
- **New tools** in the C++ Standard Library

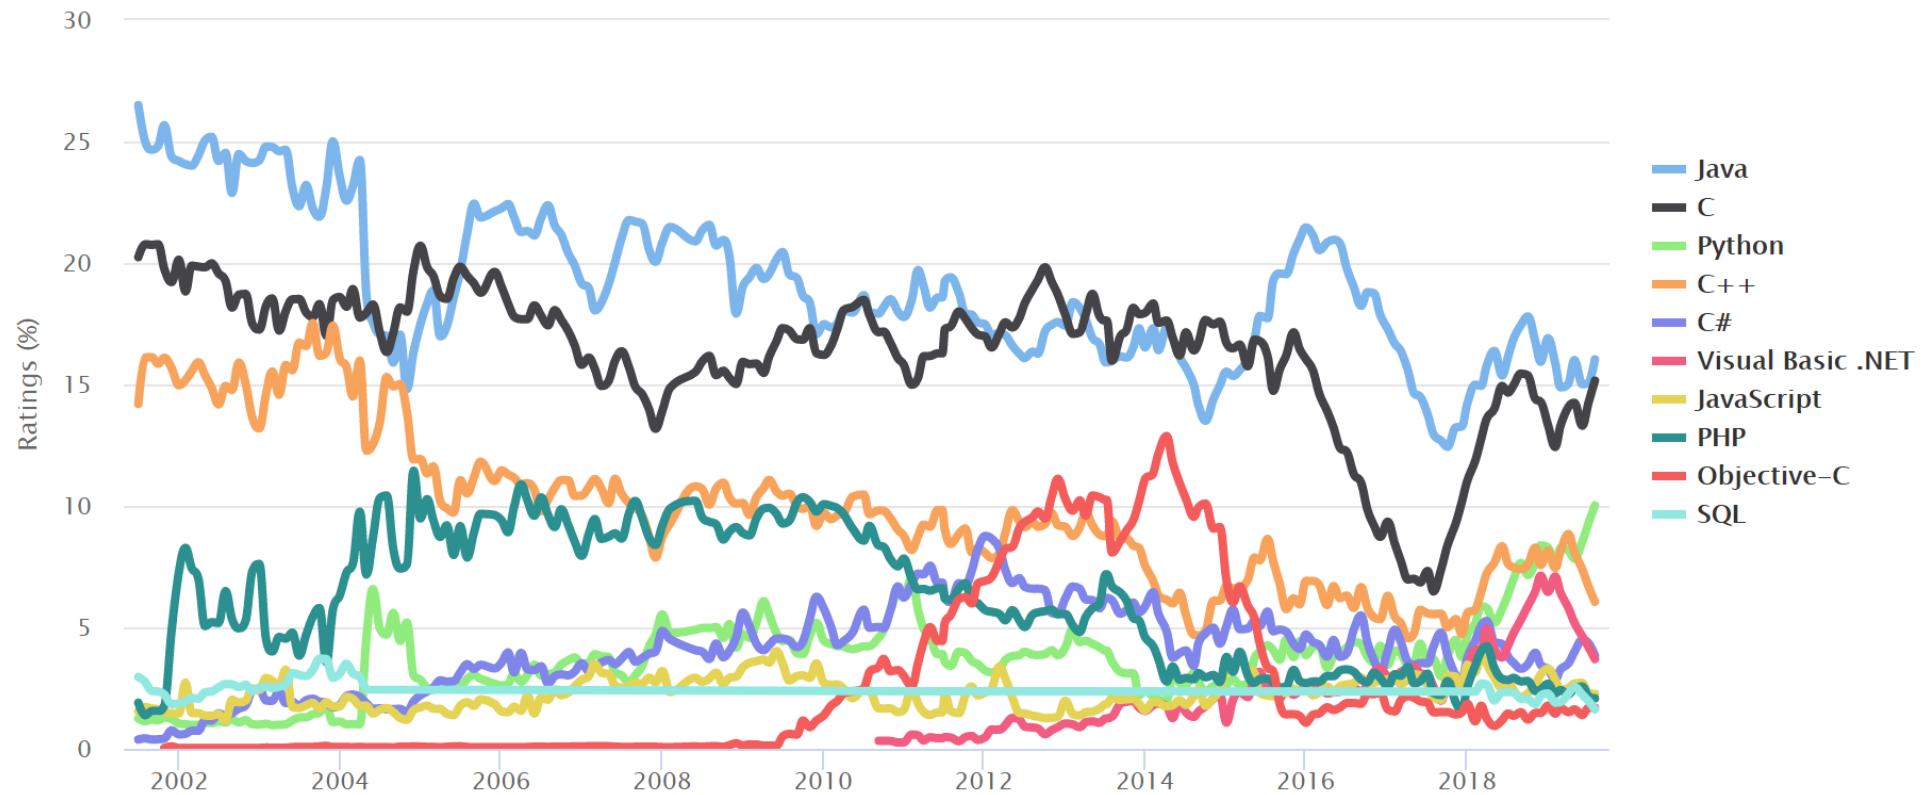
EPAM is an active participant of the C++ standardization process and contributed some performance-related features to C++20 :-)

## C++ ON THE MARKET

# Tiobe Index: August 2019

TIOBE Programming Community Index

Source: [www.tiobe.com](http://www.tiobe.com)



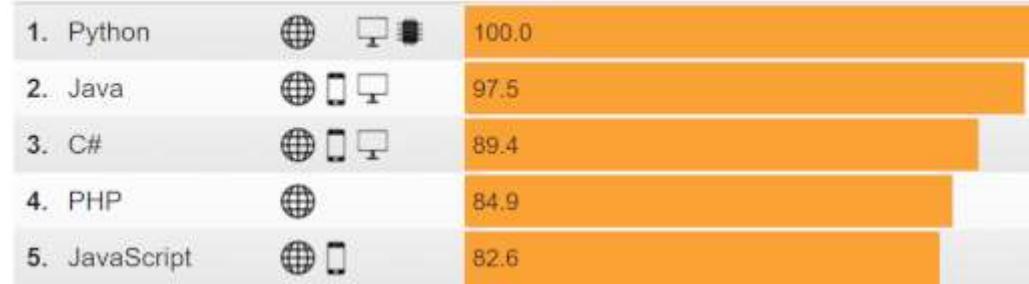
# IEEE Spectrum: The 2018 Top Programming Languages

Language Rank	Types	Spectrum Ranking
1. Python		100.0
2. C++		99.7
3. Java		97.5
4. C		96.7
5. C#		89.4
6. PHP		84.9
7. R		82.9
8. JavaScript		82.6
9. Go		76.4
10. Assembly		74.1

# IEEE Spectrum: The 2018 Top Programming Languages

## WEB

Language Rank    Types    Spectrum Ranking



## ENTERPRISE

Language Rank    Types    Spectrum Ranking



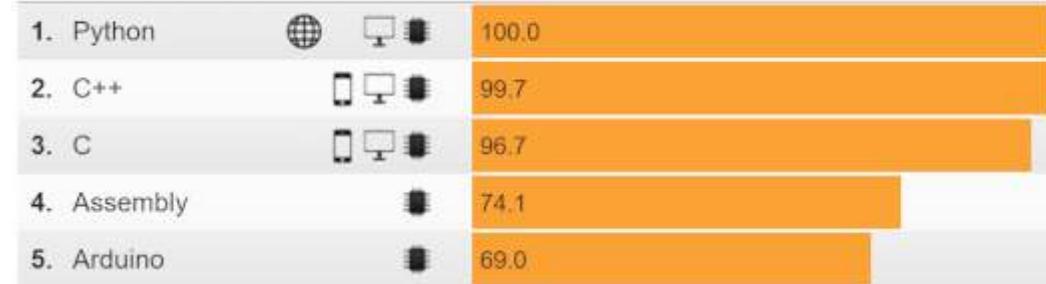
## MOBILE

Language Rank    Types    Spectrum Ranking



## EMBEDDED

Language Rank    Types    Spectrum Ranking



# Top 10 GitHub languages (# of pushes)

# Ranking	Programming Language	Percentage (Change)	Trend
1	JavaScript	27.652% (+3.768%)	
2	Python	14.915% (+0.623%)	
3	Java	11.023% (+0.831%)	
4	C++	7.068% (-0.227%)	^
5	PHP	5.718% (-1.811%)	▼
6	C#	4.453% (-1.978%)	
7	Shell	4.146% (-0.627%)	
8	Ruby	4.118% (+0.001%)	
9	Go	4.082% (-0.015%)	
10	TypeScript	3.760% (+0.510%)	^

# Facts about C++

---



<https://blog.jetbrains.com/clion/2015/07/infographics-cpp-facts-before-clion>

# Facts about C++

---

**~4.4 million C++ devs  
~1.9 million C devs**

#1

There are 4.4 million C++ developers and nearly 2 million C developers in the world.

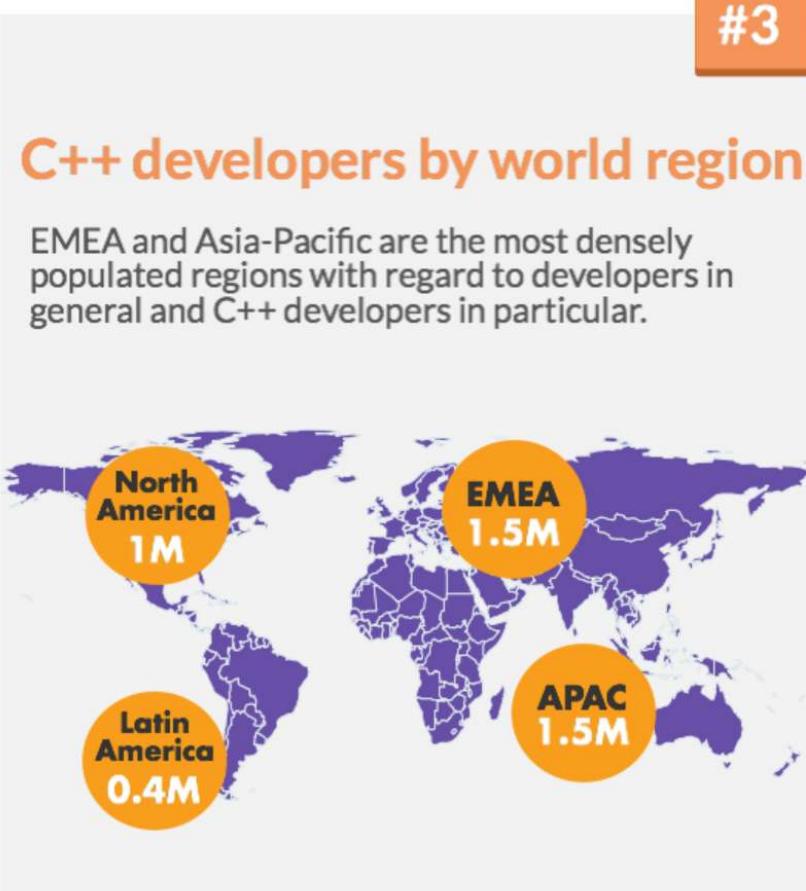


# Facts about C++

**~4.4 million C++ devs  
~1.9 million C devs**

#1

There are 4.4 million C++ developers and nearly 2 million C developers in the world.



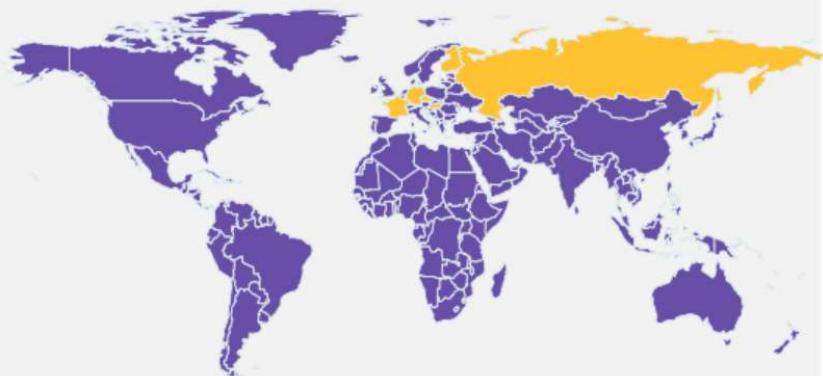
# Facts about C++

---

#4

## Where C++ is relatively ahead of other languages

C++ is relatively more popular than other languages and technologies in Russia, Czech Republic, Hungary, France, Singapore, Finland, Israel and Germany.



# Facts about C++

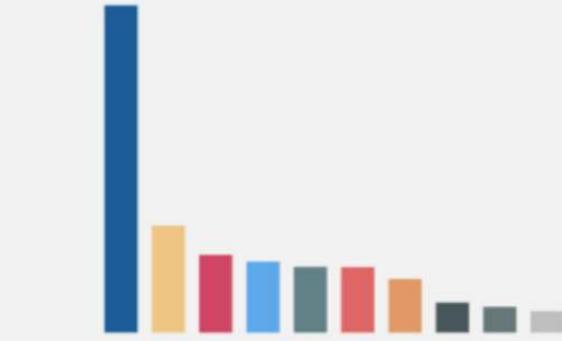
#4

## Where C++ is relatively ahead of other languages

C++ is relatively more popular than other languages and technologies in Russia, Czech Republic, Hungary, France, Singapore, Finland, Israel and Germany.



#5



■ Finance ■ Banking ■ Games  
■ Front Office ■ Telecoms ■ Electronics  
■ Investment Banking ■ Marketing  
■ Manufacturing ■ Retail

## Industry distribution

Based on the C/C++ job ads we analyzed, C++ is most used in industries like Finance and Banking.

# The Modern IT world is built on C++

---

# The Modern IT world is built on C++

---

- **Low-Level**

- Operating systems and drivers
- Embedded devices (from tiny to huge scale)
- Autonomous cars

# The Modern IT world is built on C++

---

- **Low-Level**

- Operating systems and drivers
- Embedded devices (from tiny to huge scale)
- Autonomous cars

- **High performance libraries**

- Scientific (NASA, CERN, ...)
- Hardware accelerated calculations
- AI, machine learning
- Image recognition and mapping
- Graphics
- Maps and navigation
- All kinds of frameworks (Electron, ...)

# The Modern IT world is built on C++

---

- Low-Level

- Operating systems and drivers
- Embedded devices (from tiny to huge scale)
- Autonomous cars

- High performance libraries

- Scientific (NASA, CERN, ...)
- Hardware accelerated calculations
- AI, machine learning
- Image recognition and mapping
- Graphics
- Maps and navigation
- All kinds of frameworks (Electron, ...)

- High performance servers

- Stock Exchanges
- High Frequency Trading
- Data centers (Google, Facebook, Amazon, ...)
- Databases (MySQL, ...)
- Telecom

# The Modern IT world is built on C++

---

- **Low-Level**

- Operating systems and drivers
- Embedded devices (from tiny to huge scale)
- Autonomous cars

- **High performance libraries**

- Scientific (NASA, CERN, ...)
- Hardware accelerated calculations
- AI, machine learning
- Image recognition and mapping
- Graphics
- Maps and navigation
- All kinds of frameworks (Electron, ...)

- **High performance servers**

- Stock Exchanges
- High Frequency Trading
- Data centers (Google, Facebook, Amazon, ...)
- Databases (MySQL, ...)
- Telecom

- **End-user applications**

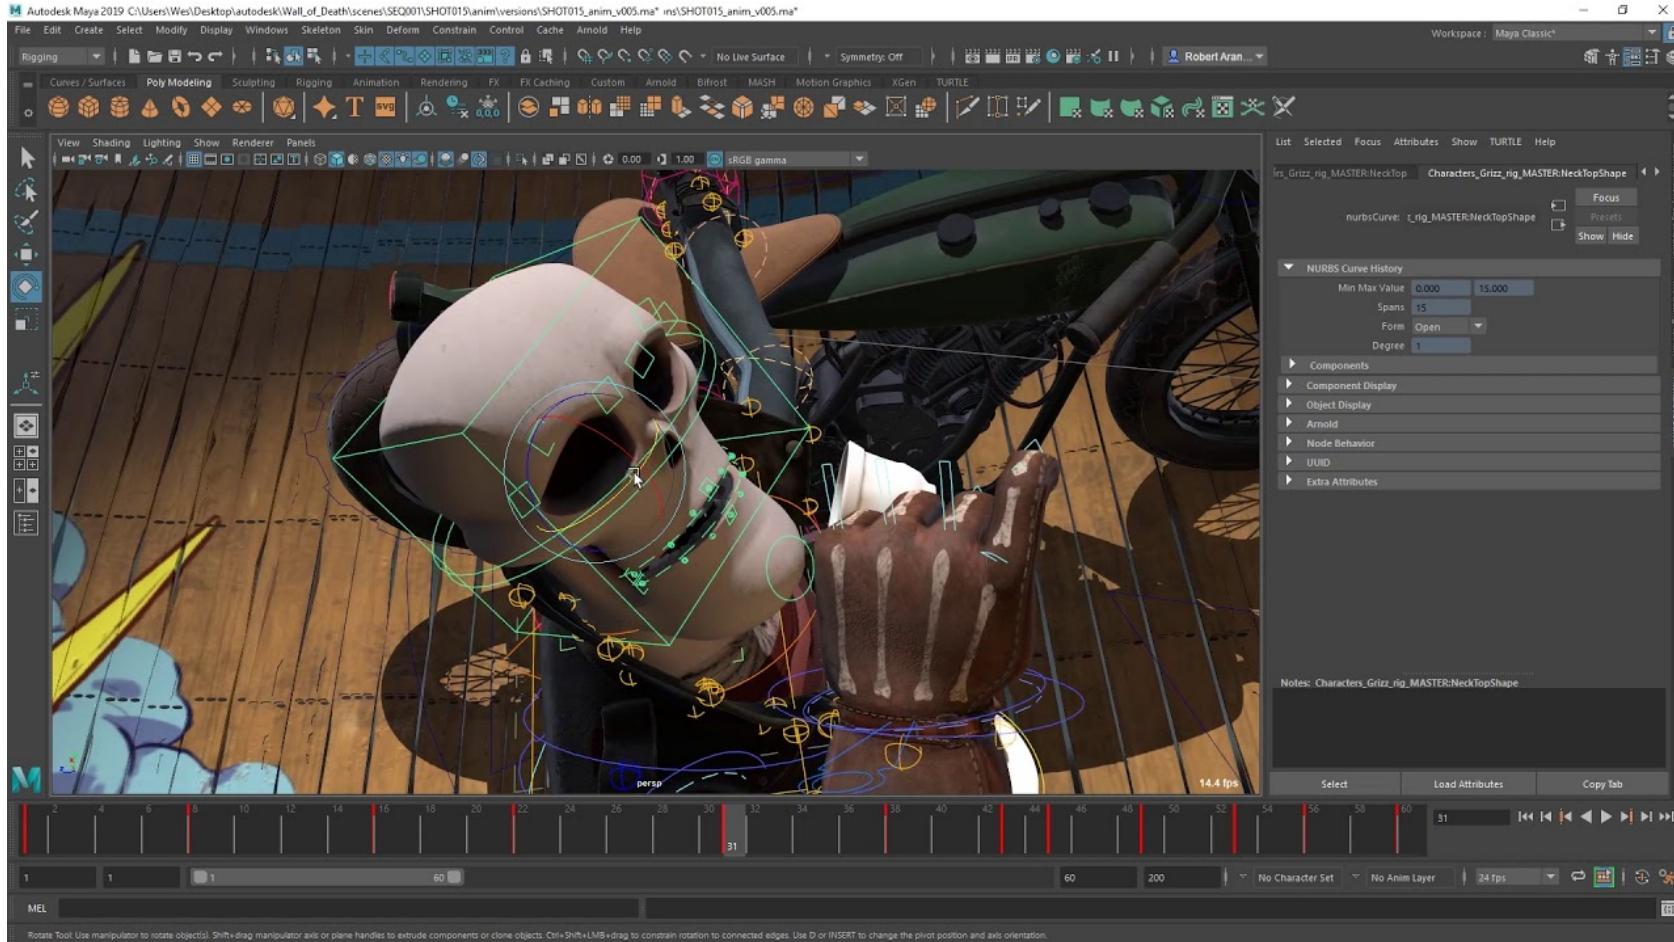
- Games
- Powerful modeling and simulation software (Adobe, Maya, Autodesk, Pixar, ...)
- Compilers

# Windows

---



# 3D Computer Animation

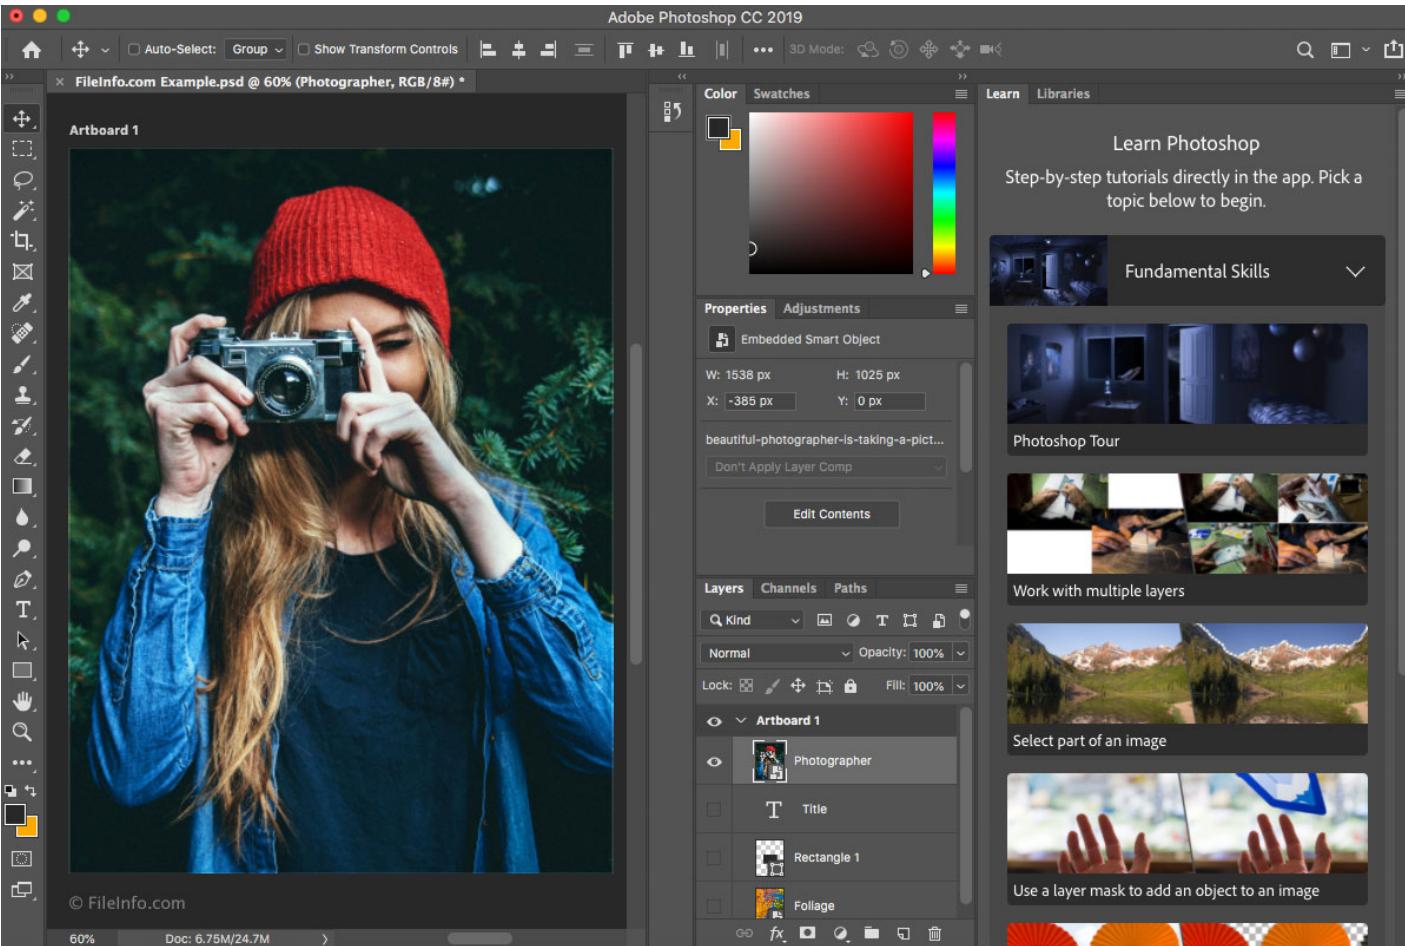


# Movies

---



# Raster graphics editors



# Light Projections

---



# Games

---



# C++ is the most efficient programming language in the world

---

# C++ is the most efficient programming language in the world

---

1

Learn Modern C++ and stay current with a fast evolving language

- follow and *learn C++ releases*
- follow *C++ conferences and blogs*
- hire *C++ trainers*

# C++ is the most efficient programming language in the world

---

## 1 Learn Modern C++ and stay current with a fast evolving language

- follow and *learn C++ releases*
- follow *C++ conferences and blogs*
- hire *C++ trainers*

## 2 Use C++ in your projects to

- *decrease money expenditure* on computational resources
- boost *performance*
- limit *jitter* (improve response time)
- limit *power usage*
- limit *memory footprint*
- have *more control over hardware*

